



Google BigTable

DS Seminar



23-05-2016

Agenda

- A Short history of Google Motivation
- What is GFS ?
- What is Chubby ?
- What is Map Reduce ?
- What is BigTable ?
 - Data Model
 - API
 - Building Blocks
 - Implementation

Google Motivation

- ❑ Undoubtedly, Google has **a lot of data**.

- ❑ Scale of data is **too large**

Even for commercial databases.

Even though **Google** is *best known* for it's **reliable** and **fast services**, but what's there working behind the scene?



Google Motivation

- ❑ Undoubtedly, there are number of aspects that matter behind this
(like **Hardware, Software, Operating System, Best Staff** in the world etc.)
- ❑ But, What I am going to explain here is the Software part.

- ❑ GFS
- ❑ Chubby
- ❑ Map Reduce
- ❑ BigTable



What is GFS ?

- ❑ **GFS** stands for **Google File System**.
- ❑ It's a **Proprietary**(means for their personal use, not open source) **distributed file system** developed by Google for their services.
- ❑ It is specially designed to provide **efficient, reliable** access to data using **large clusters** of **commodity hardware**, means they are using low cost hardware, not state-of-the-art computers. Google uses relatively inexpensive computers running Linux Operating System and the GFS works just fine with them !

What is Chubby ?

- ❑ Chubby is a **Lock Service**. (It's related to gain access of Shared resources)
- ❑ It is used to **synchronize accesses** to **shared resources**.
- ❑ It is now used as a replacement of Google's Domain Name System.



What is Map Reduce ?

- ❑ MapReduce is a software framework that process **massive** amounts of unstructured data.
- ❑ It allows developers to write programs that process data in **parallel** across a distributed cluster of processors or stand-alone computers.
- ❑ It is now used by Google mainly for their **Web Indexing** Service, applied since 2004.
- ❑ **Map()** procedure performs all the process related to **Filtering** and **Sorting**.
- ❑ **Reduce()** procedure performs all the **Summary** related operations.

What is BigTable?

- ❑ BigTable is a **distributed storage system** for **managing structured data** built on Google File System, Chubby Lock Service, SSTable (log-structured storage like LevelDB) and a few other Google technologies.
- ❑ **Designed to Scale** to a very large size: petabytes of data across thousands of commodity servers
- ❑ Most important point, It's a **Non-Relational Database**.
- ❑ It uses amazing **Load Balancing Structure** so that it runs on Commodity Hardware.
- ❑ It uses **Snappy** compression utility for compacting the data.

What is BigTable?

- ❑ Distributed
- ❑ Column – Oriented
- ❑ Multidimensional
- ❑ High Availability
- ❑ High Performance
- ❑ Store System
- ❑ Self-managing



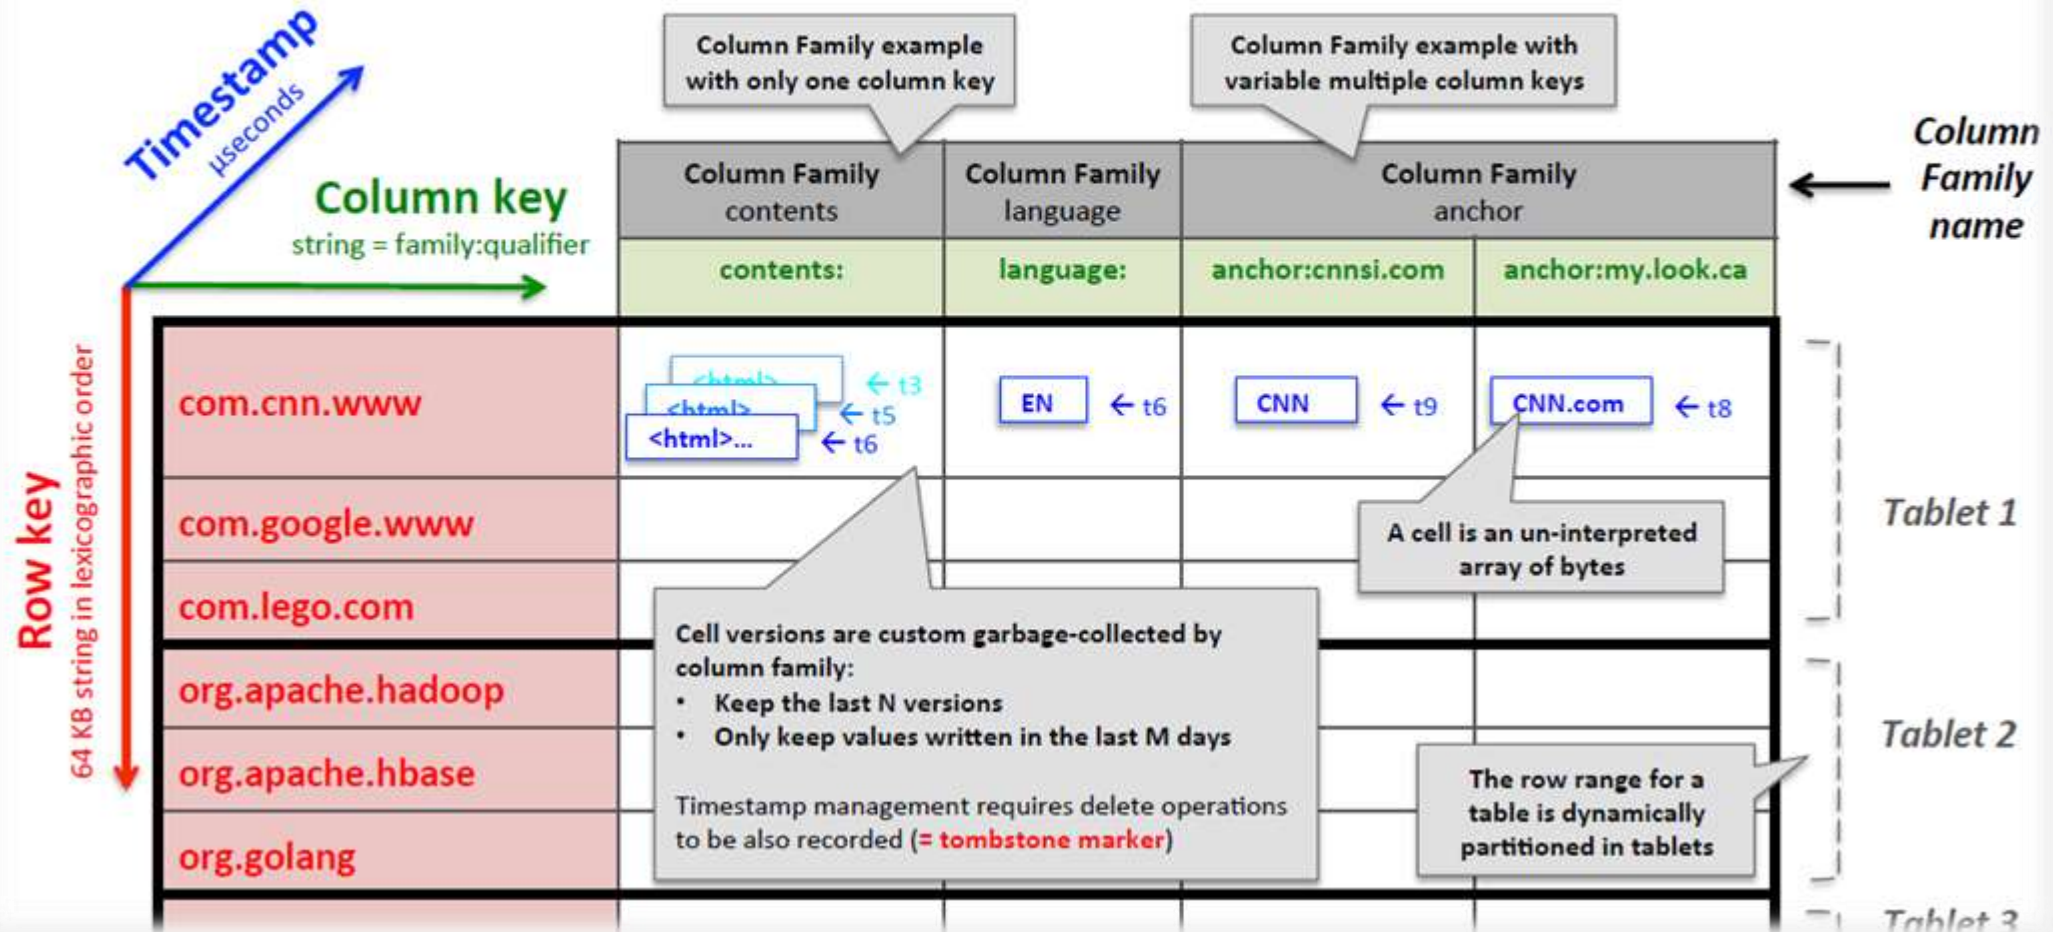
What is BigTable?

- Used by more than 60 Google products
 - Google Analytics
 - Google Finance
 - Personalized Search
 - Google Documents
 - Google Earth
 - Google Fusion Tables
 - ...
- Used for variety of demanding workloads
 - Throughput oriented batch processing
 - Latency sensitive data serving



Data Model

Data model



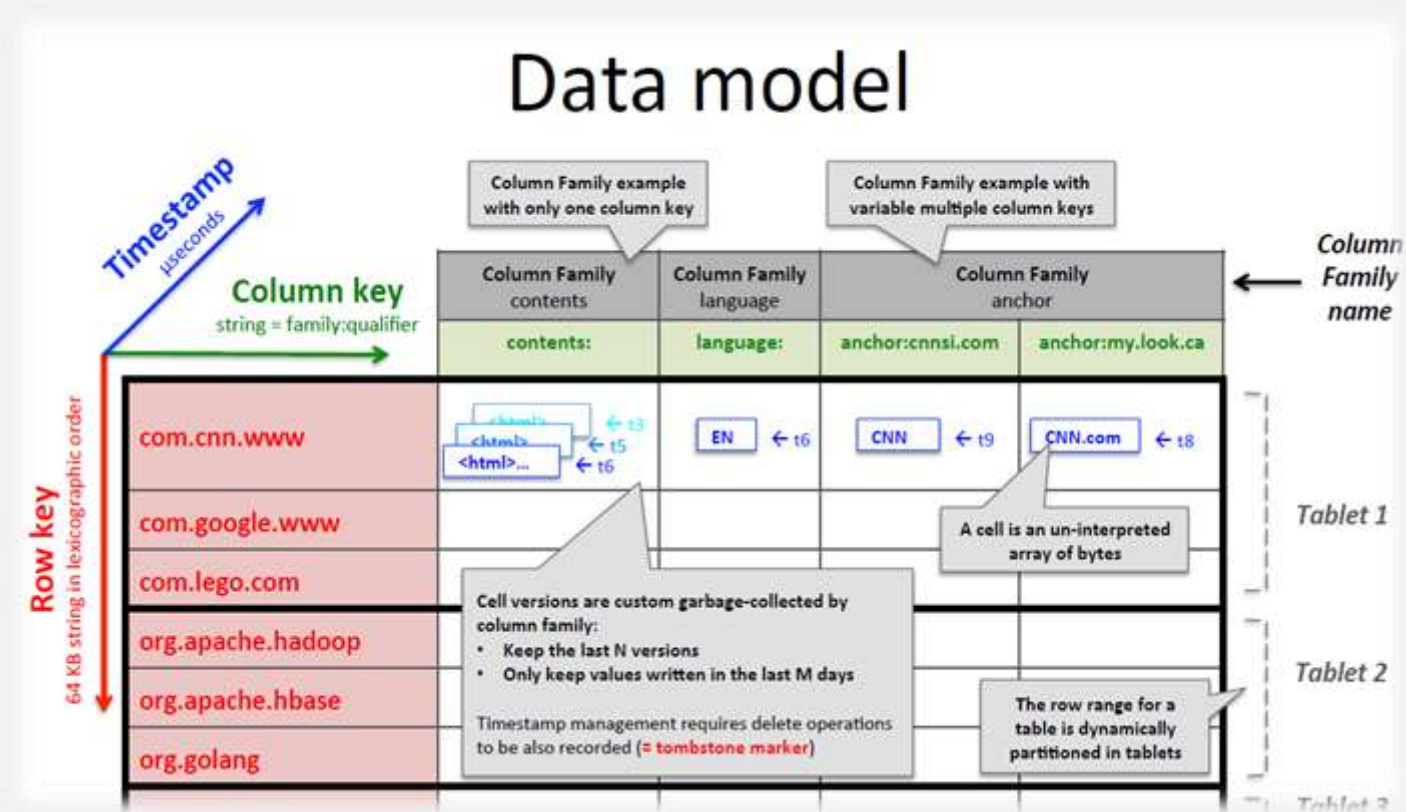
Data Model

❑ Each table is a **Multi-Dimensional Sparse Map** (Memory Efficient hash-map implementation).

❑ The table consists of

- (1) Rows,
- (2) Columns
- (3) Each cell has a Time Version
(Time-Stamp).

❑ Time Version results in multiple copies of each cell with different times,

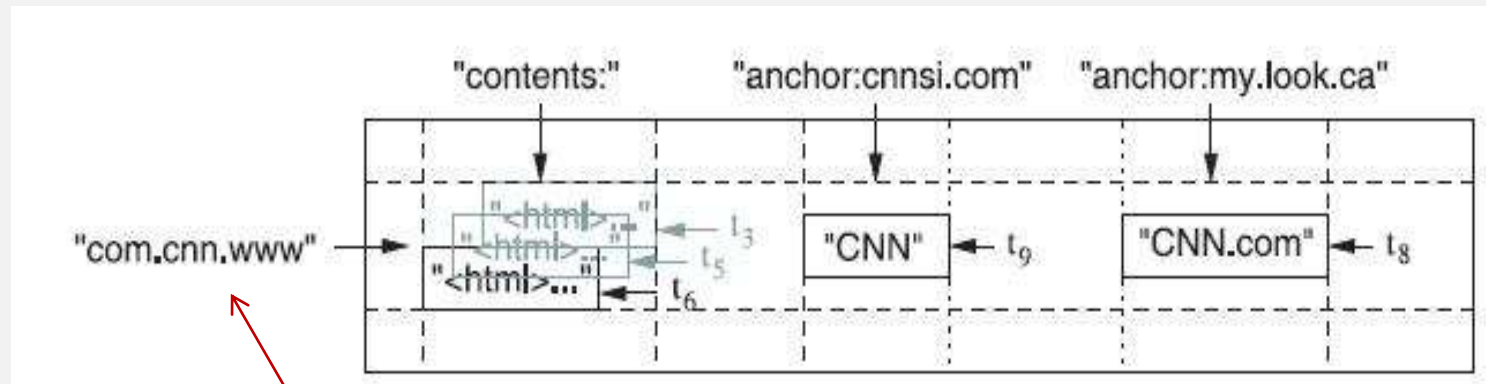


Data Model

- ❑ Time Version results in multiple copies of each cell with different times, resulting **Unimaginable Redundancy** which is requirement for Google services, so don't ever think it as a drawback of this system.
- ❑ Google does Web Indexing to get the data of all the websites. They store all the URLs, their titles, time-stamp and many more required fields
- ❑ Web Indexing :- indexing the contents of a website

Data Model-Row

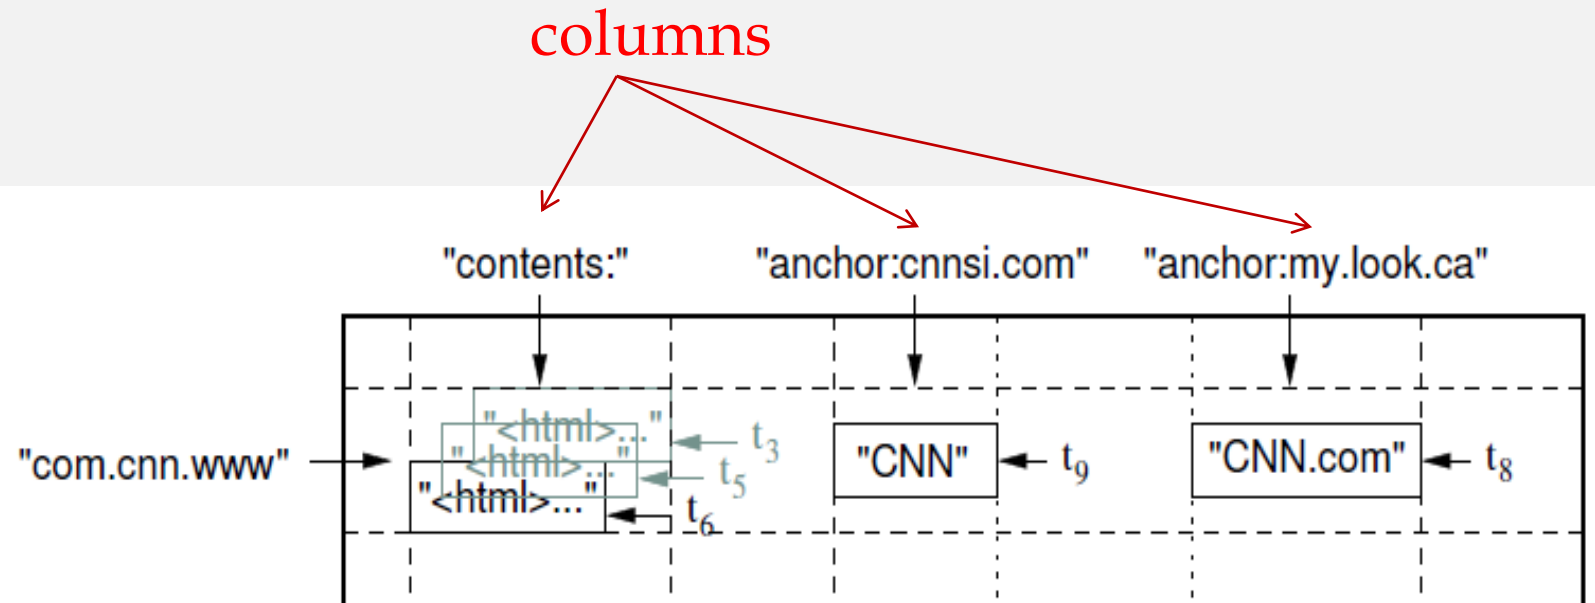
- ❑ The row keys in a table are **arbitrary strings**.
- ❑ Data is maintained in **lexicographic order** by row key
- ❑ Each **row range** is called a **tablet**, which is the unit of distribution and load balancing.



row

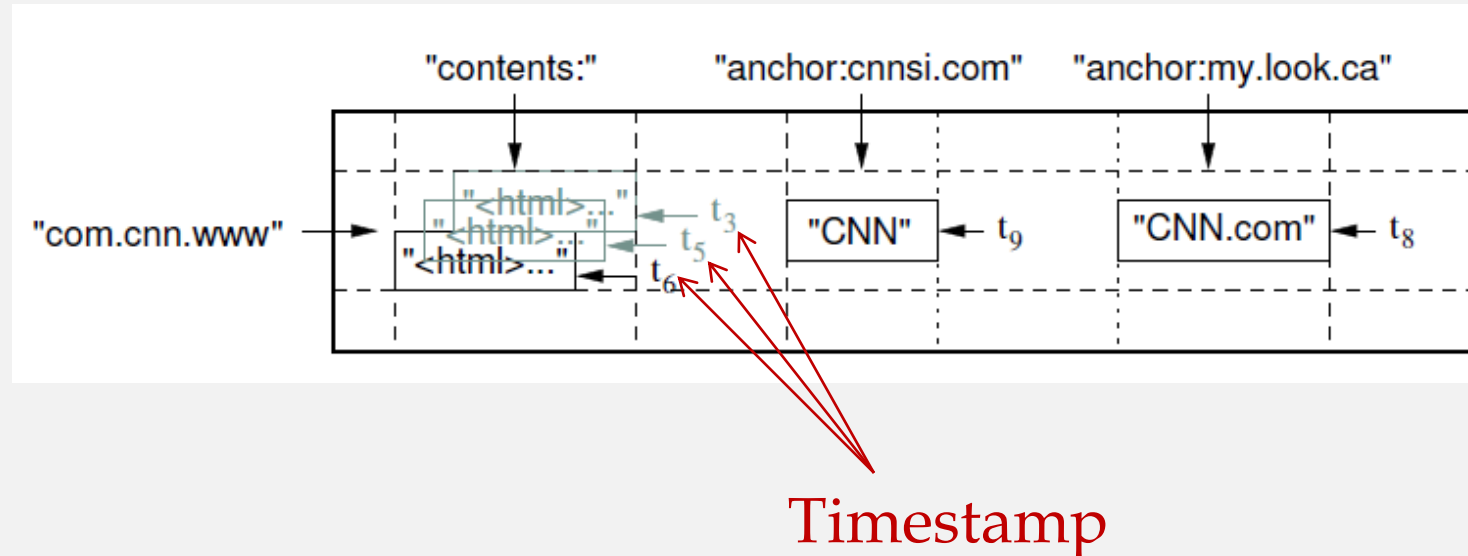
Data Model-Column

- ❑ Column keys are grouped into sets called **column families**.
- ❑ Data stored in a column family is usually of the same type
- ❑ A column key is named using the syntax: family : qualifier.
- ❑ Column family names must be printable , but qualifiers may be arbitrary strings.



Data Model-Timestamp

- ❑ Each cell in a Bigtable can contain multiple versions of the same data
- ❑ Versions are indexed by **64-bit integer** timestamps
- ❑ Timestamps can be assigned:
 - automatically by Bigtable , or
 - explicitly by client applications



Data Model

key

value

Key : row key, column family, column, time stamp

row key

column family

column

time stamp

value

row key = byte array , diset sbg primary key
data diurut (disortir) berdasarkan row key

row key

column family

column

time stamp

value

row1

column family1

column1

t1

'Java'

row1

column family1

column2

t1

'C'

row1

column family2

column1

t1

'C++'

row2

column family1

column1

t1

'Fortran'

row2

column family1

column2

t1

'Perl'

API

❑ The Bigtable API provides functions :

- Creating and deleting tables and column families.
- Changing cluster , table and column family metadata.
- Support for single row transactions
- Allows cells to be used as integer counters
- Client supplied scripts can be executed in the address space of servers
- Input and output of Map/Reduce jobs

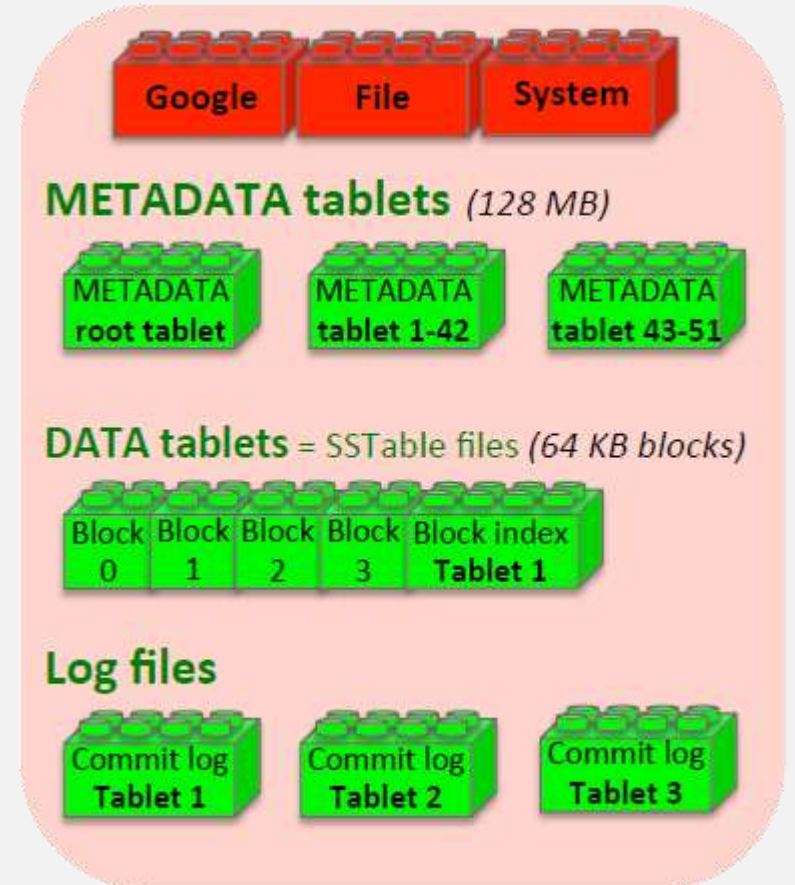
API

Bigtable API

- Tablet server `AddServer(tabletServer) / RemoveServer(tabletServer)`
- Table `CreateTable(table) / DeleteTable(table)`
- Column family `CreateColumnFamily(columnFamily) / DeleteColumnFamily(columnFamily)`
- Table access control rights and metadata `SetTableFlag(table, flags) / . . .`
- Column family access control rights and metadata `SetColumnFamilyFlag(table, colfamily, flags) / . . .`
- Cell value `Put(rowkey, columnkey, value) / Get(rowkey, columnkey) / Delete(rowkey, columnkey)`
- Look up value from individual row `Has(rowkey, columnfamily) / . . .`
- Look up values from table (=MapReduce like RPC) `Scan(rowFilter, columnFilter, timestampFilter)`
 - Can iterate over multiple column families
 - Can limiting rows/columns/timestamps
- **Single-row transactions (atomic read-modify-write sequence)**
- **No support for general transactions across row keys**
- Cells can be used as integer counters `Increment(rowkey, columnkey, increment)`
- Execution of read-only client-supplied scripts in the address spaces of the servers: *Sawzall*
 - <http://research.google.com/archive/sawzall.html>
- Bigtable can be used with MapReduce (for input and/or output)

Building Blocks

- ❑ BigTable is composed of several other innovative, distribution oriented components.
- ❑ Google File System (GFS)
- ❑ SSTable
- ❑ Chubby



Building Blocks

- ❑ Google File System (GFS) :

Used to store log and data files

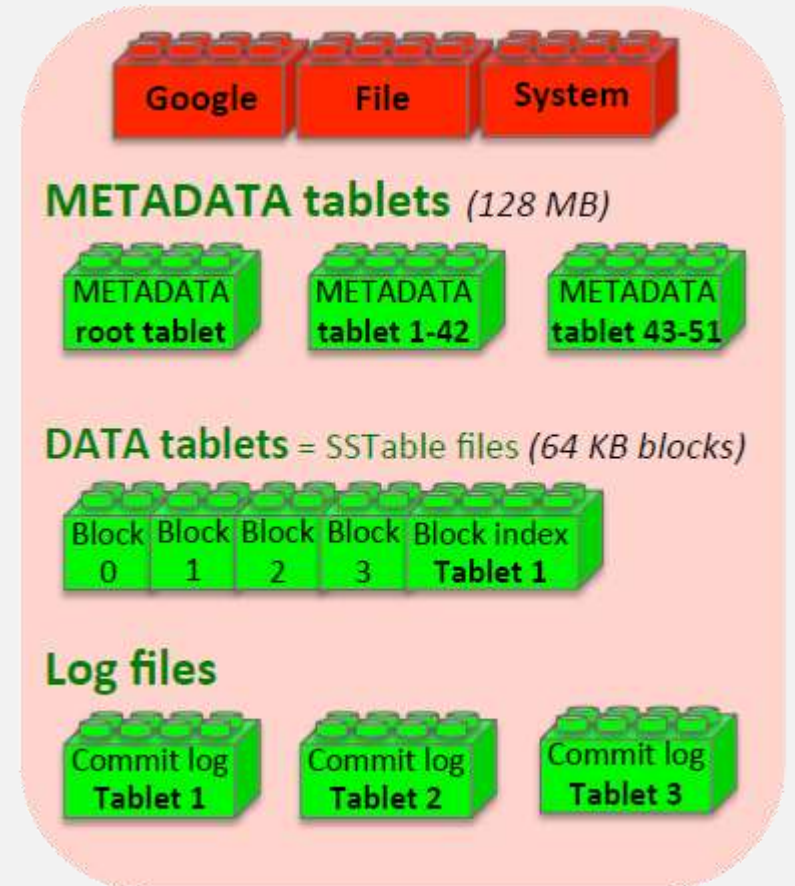
- ❑ SSTable (Sorted String Table) :

Used to store table data in GFS

used to **store** and **retrieve** the pairs <Key, Value>

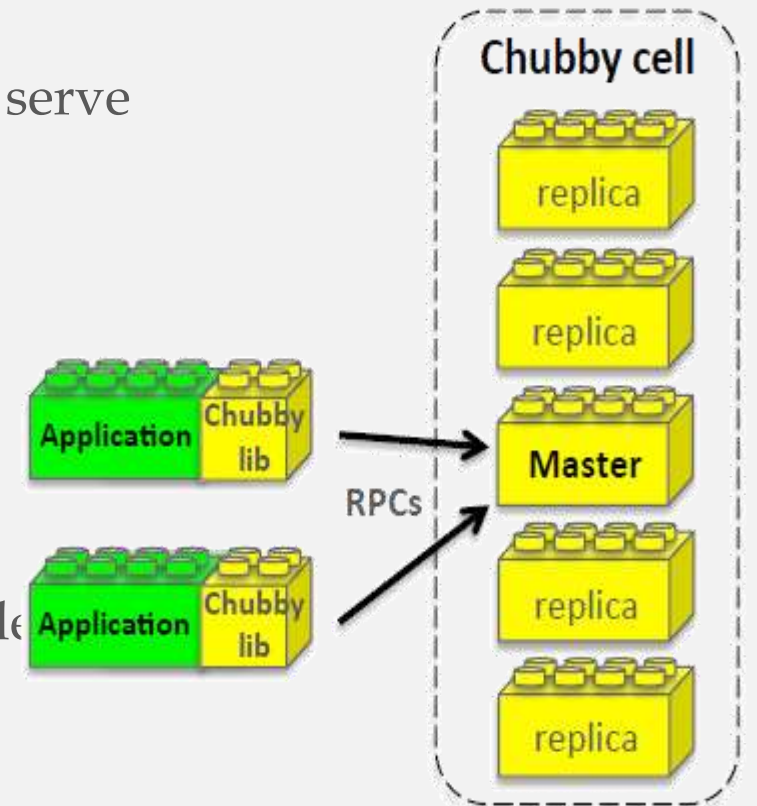
Used as **pointers** to pairs <Key, Value>

Stored in GFS



Building Blocks

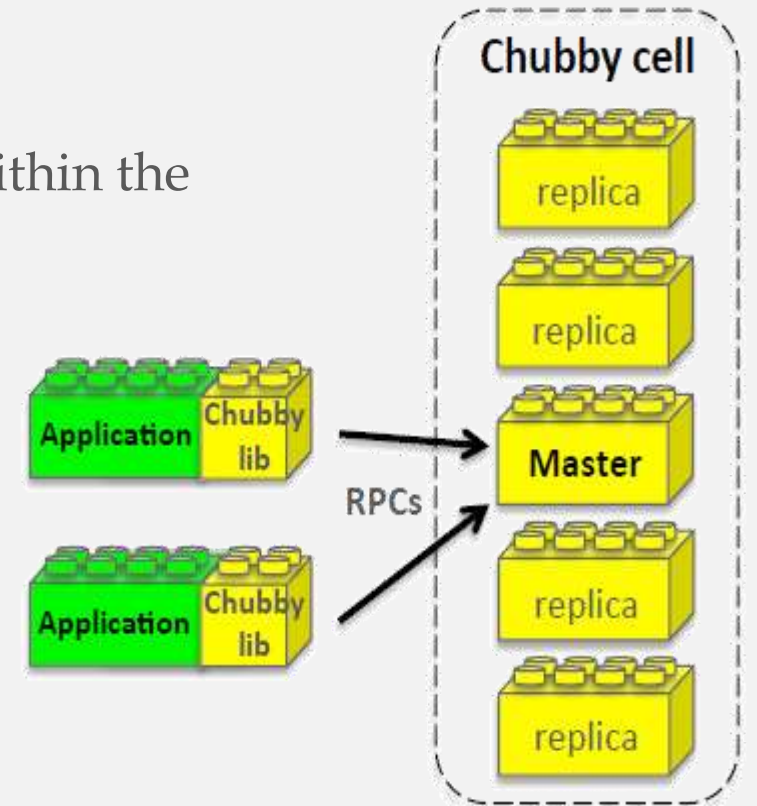
- ❑ Google Chubby :
- Chubby is High available and persistent distributes service
- Chubby service consists of **5 active replicas** with one master to serve requests
- Each directory or file can be used as a lock
- Reads and writes to a file are atomic
- Chubby client library provides consistent caching of Chubby file
- Each Chubby client maintains a session with a Chubby service



Building Blocks

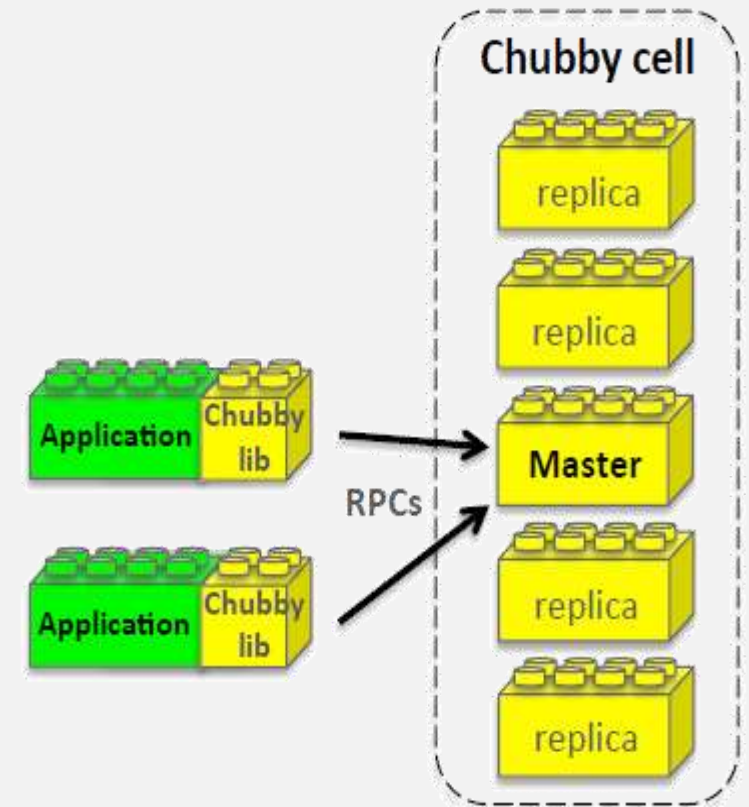
□ Google Chubby :

- Client's session expires if is unable to renew its session lease within the lease expiration time
- When a client's session expires, it loses any locks and open handles
- Chubby clients can also register callbacks on Chubby files and directories for notification of changes or session expiration




Building Blocks

- ❑ BigTable uses Chubby for a variety of tasks
 - To ensure there is at most one active master at any time
 - To store the bootstrap location of BigTable data (Root tablet)
 - To discover tablet servers and finalize tablet server deaths
 - To store BigTable schema information (column family information for each table)
 - To store access control lists (ACL)
 - Chubby unavailable = BigTable unavailable



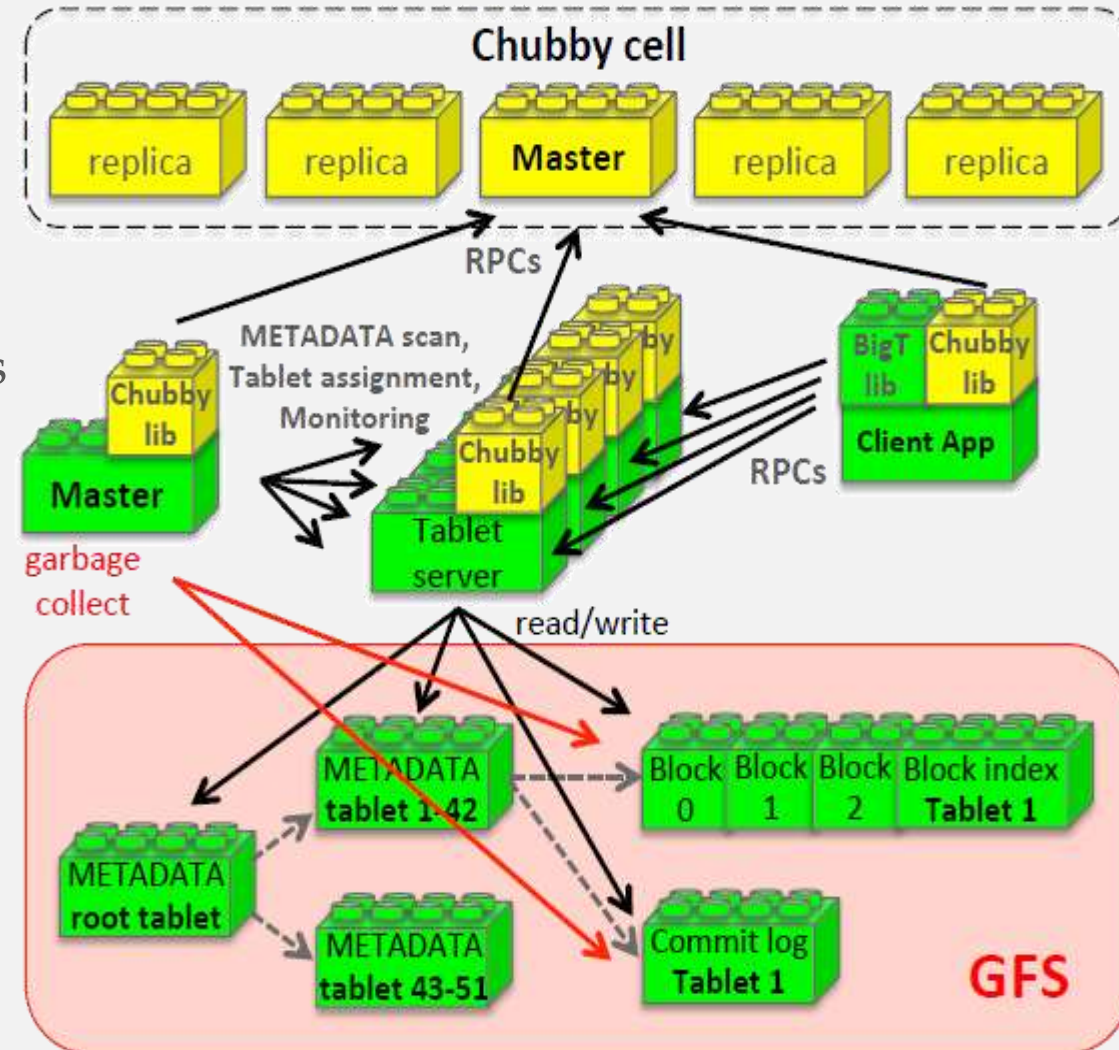
Implementation

- ❑ The implementation has three major components
 - **One** Master server
 - **Many** tablet servers
 - A **library** that is linked into every client
- ❑ BigTable runs over Google File System
- ❑ BigTable is store in a structure called SSTable. Each SSTable is divided into 64KB blocks. A SSTable can be loaded to Memory



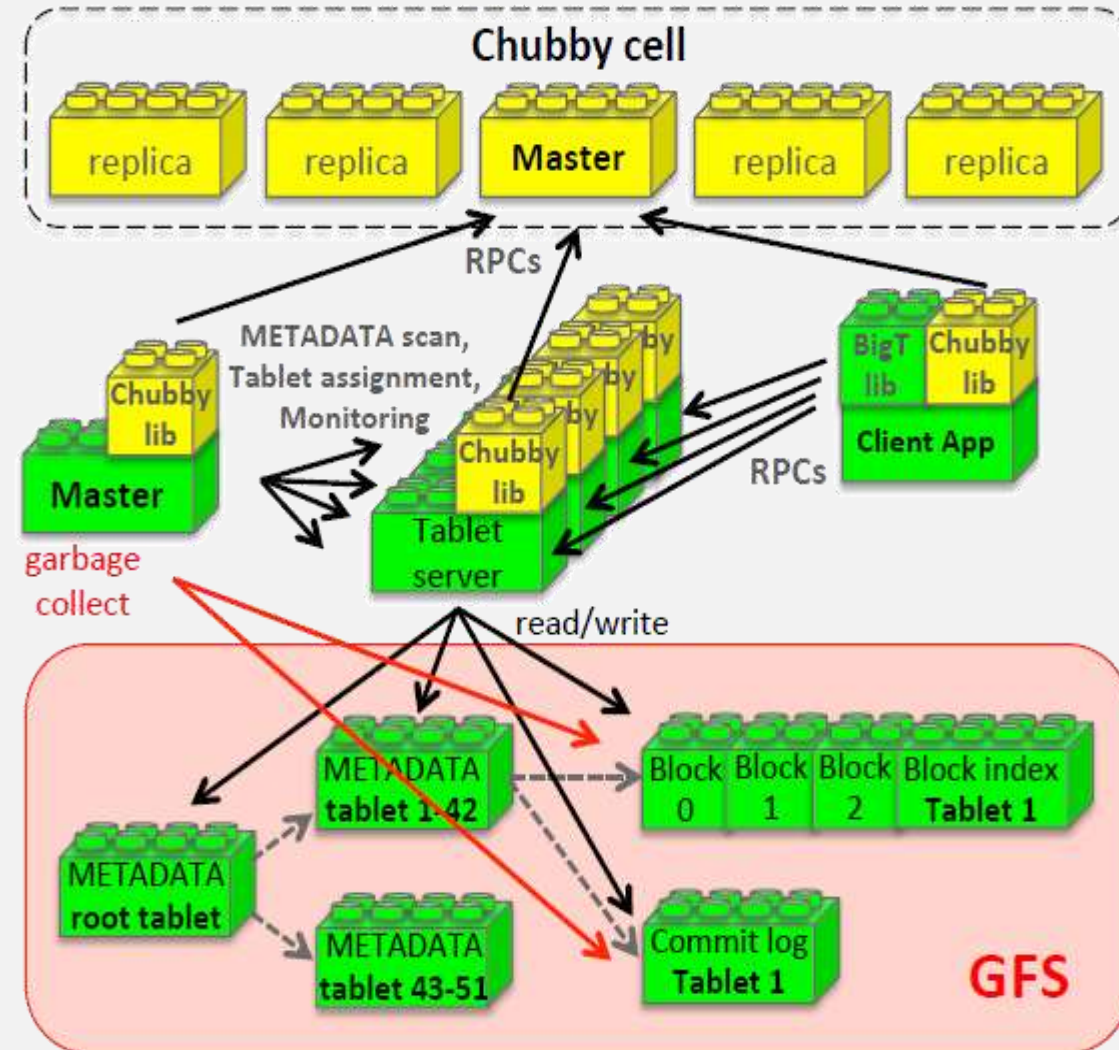
Implementation

- **One** Master server
- ✓ Assigning tablets to tablet servers
- ✓ Detecting the addition and expiration of tablet servers
- ✓ Balancing tablet server load
- ✓ Garbage collecting of files in GFS
- ✓ Handling schema changes (table creation, column family creation/deletion)




Implementation

- **Many** tablet servers
- ✓ Manages a set of tablets
- ✓ Handles read and write request to the tablets
- ✓ Splits tablets that have grown too large (100--200 MB)



Implementation

- A **library** that is linked into every client
- ✓ Do not rely on the master for tablet location information
- ✓ Communicates directly with tablet servers for reads and writes



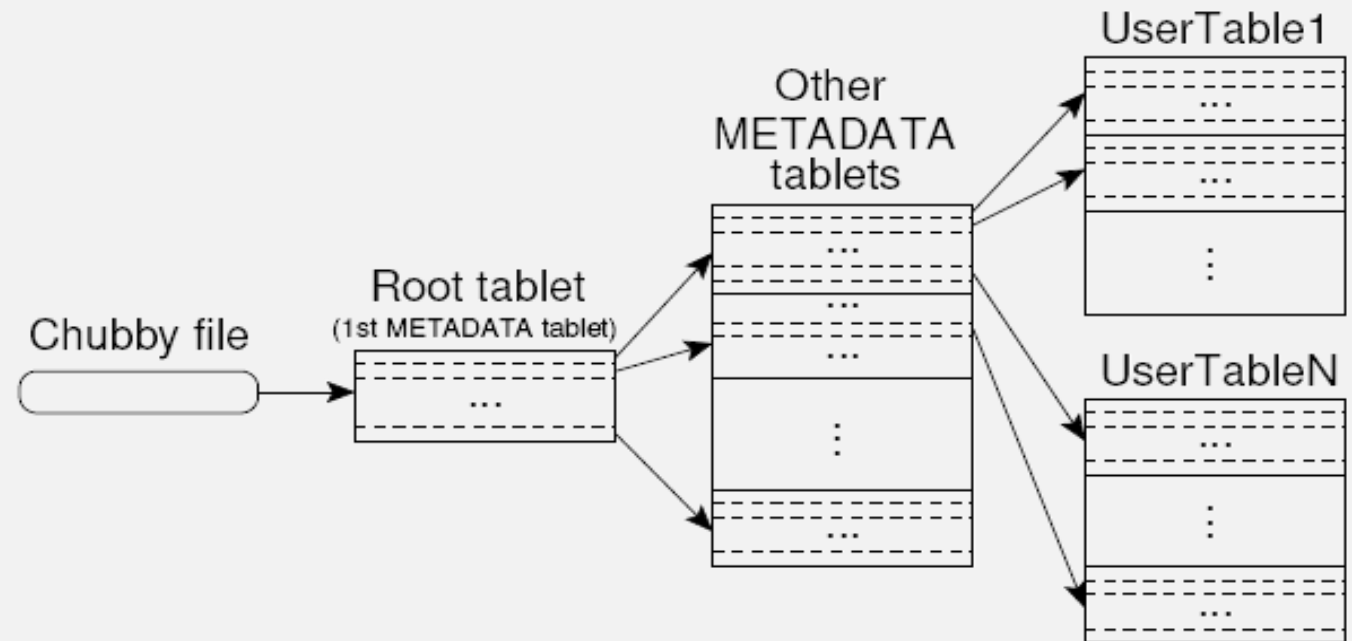
Tablet Location

❑ **Chubby File:** Provides an **namespace** to **access** the **root table**.

This is the first entry point to locate a user table.

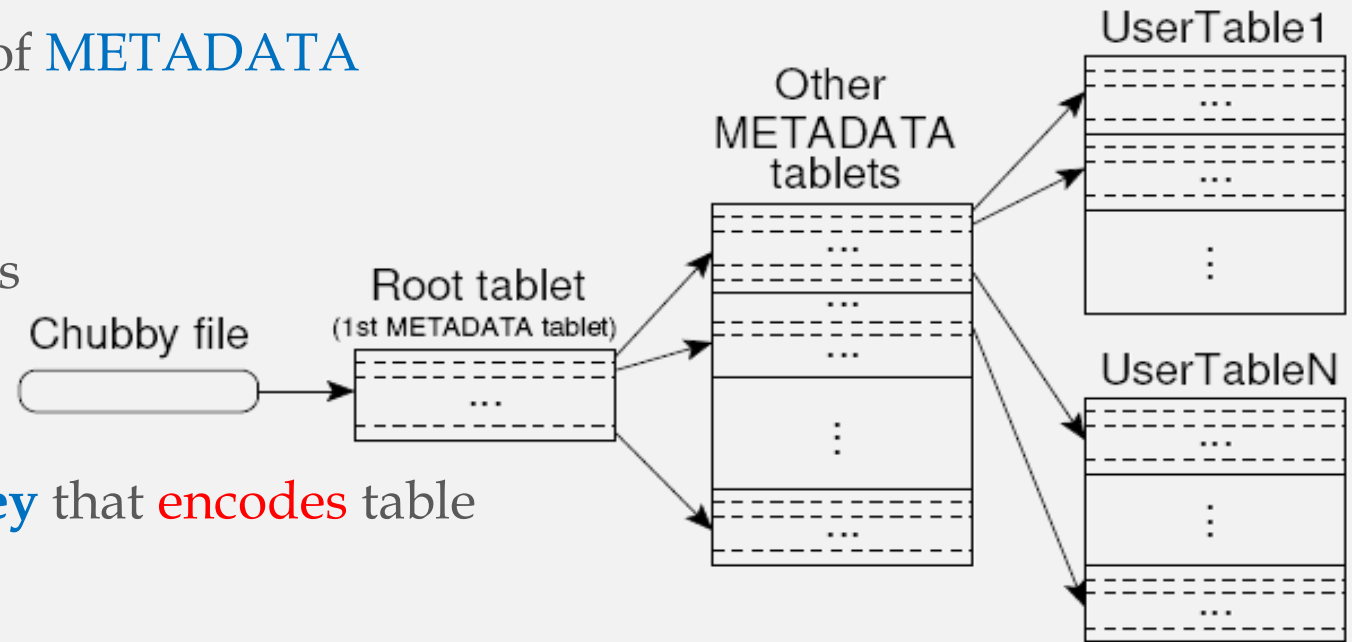
The service is distributed. The chubby service is used for:

- ✓ Bootstrap the location of BigTable
- ✓ Discover server tablets
- ✓ Finalize tablets servers deaths



Tablet Location

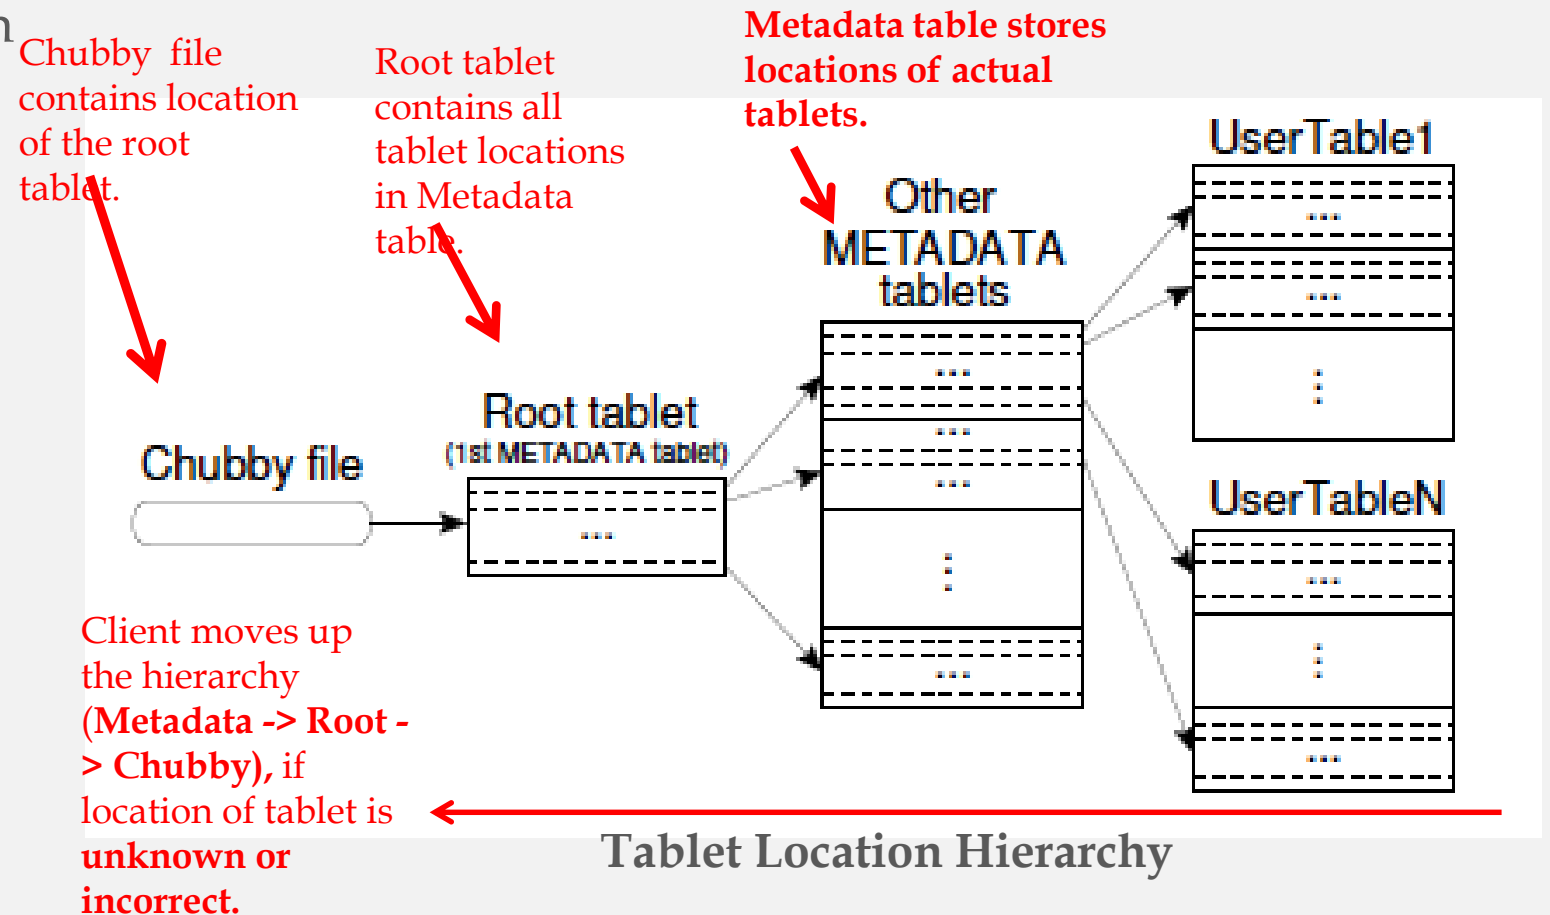
- ❑ Three level hierarchy
- ❑ Level 1: **Chubby file** containing location of the **root tablet**
- ❑ Level 2: **Root tablet** contains the location of **METADATA tablets**
- ❑ Level 3: Each **METADATA tablet** contains the location of **user tablets**
- ❑ Location of **tablet** is stored under a **row key** that **encodes** table identifier and its end row



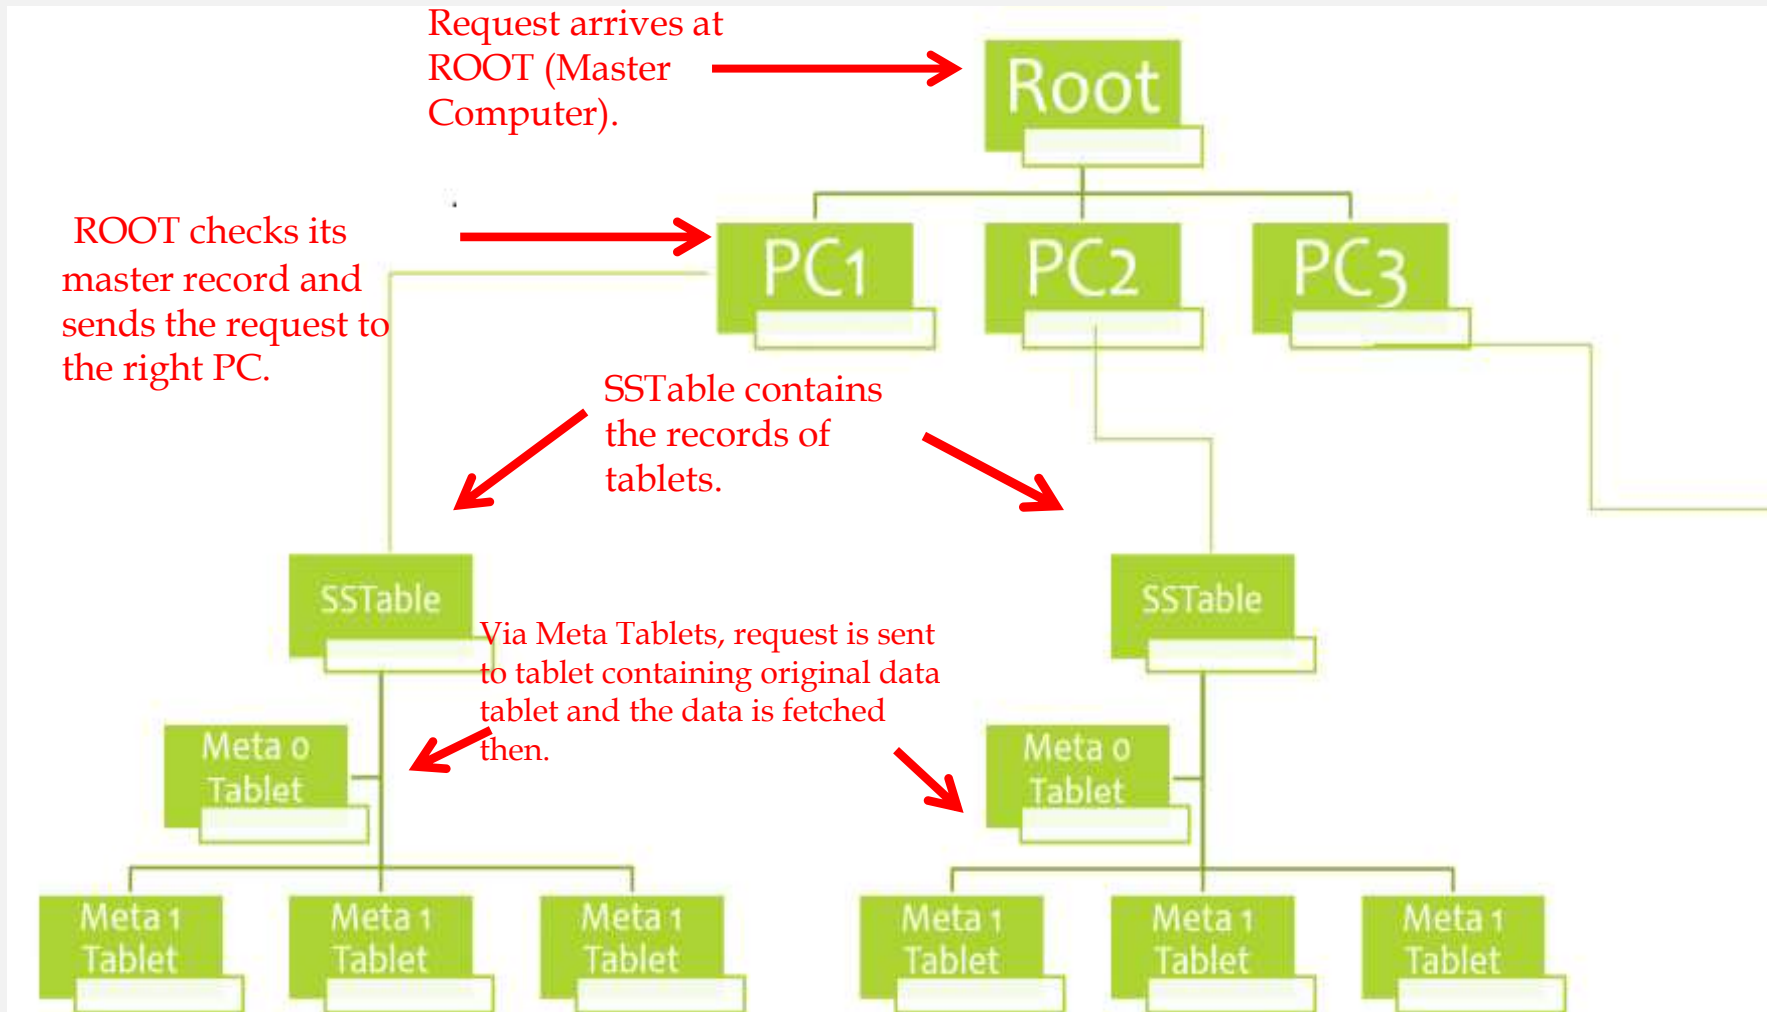
Tablet LookUp

❑ Lookup is a **three-level system**.

❑ **Benefit :- NO Big Bottleneck** in the system and it also make heavy use of **Pre-Fetching** and **Caching**



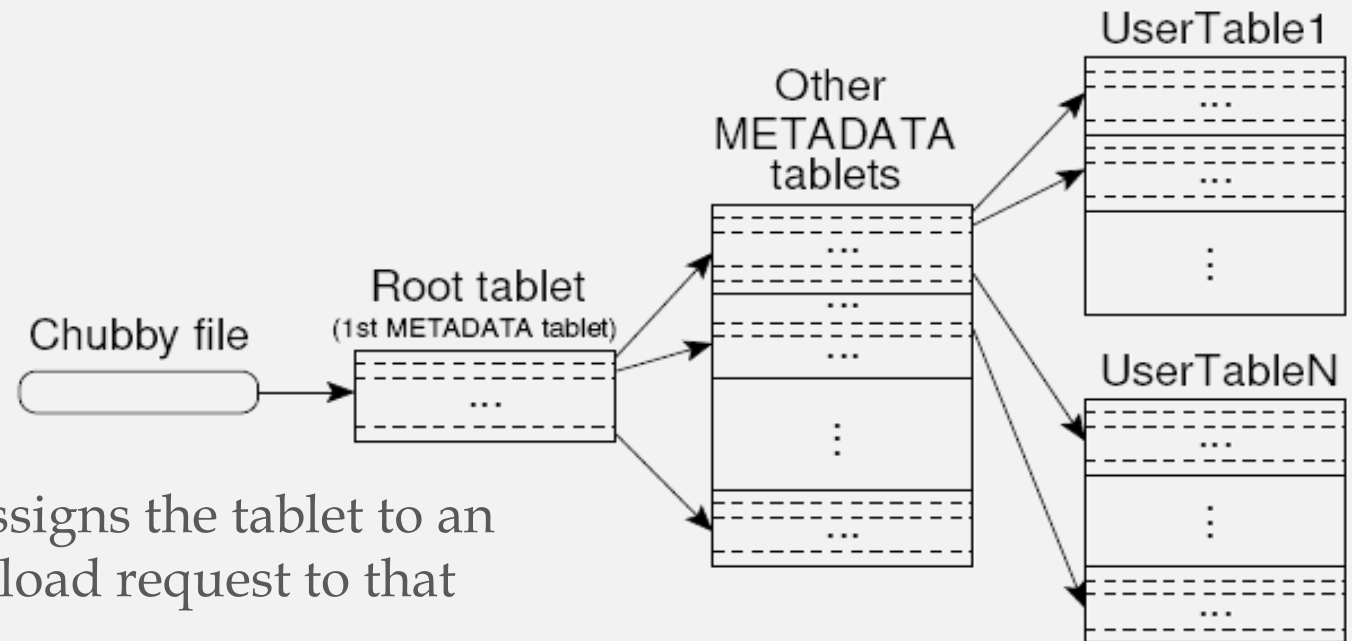
Actual Hierarchical Load Balancing Structure



This is how, it works

Tablet Assignment

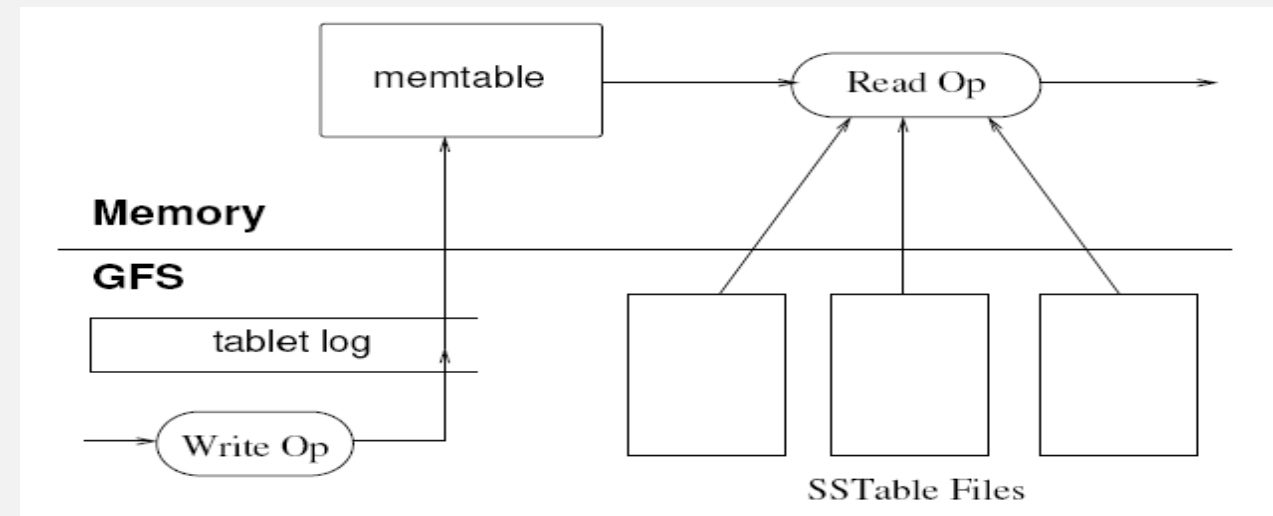
- ❑ Each tablet is assigned to **one** tablet server at a **time**
- ❑ **Master** keeps tracks of
 - the set of **live tablet servers** (tracking via Chubby)
 - the current **assignment of tablet** to tablet servers
 - the current **unassigned tablets**



- ❑ When a tablet is unassigned, the **master** assigns the tablet to an available tablet server by sending a tablet load request to that tablet server

Tablet Serving

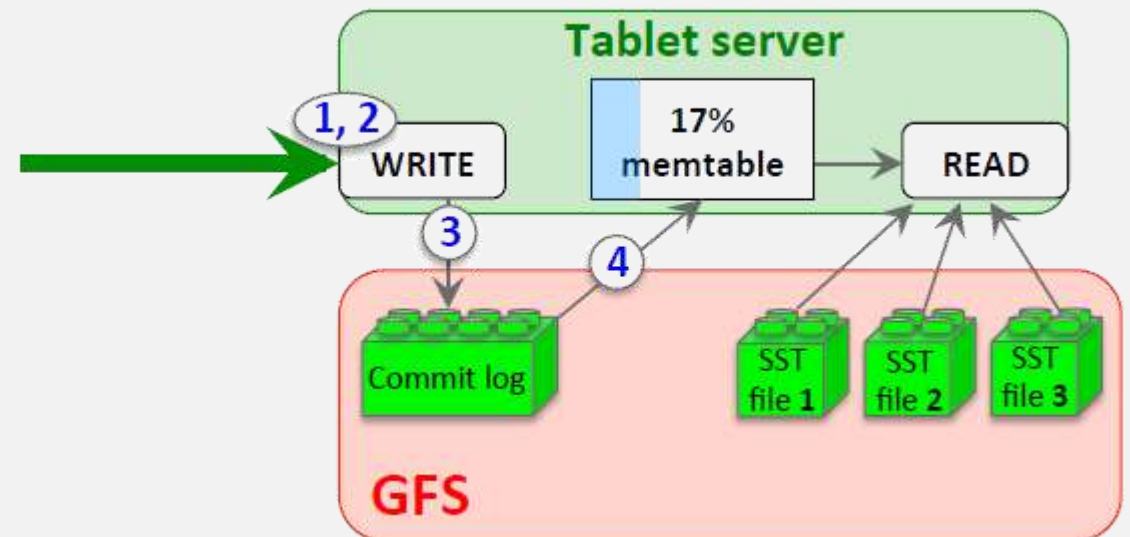
- ❑ Updates committed to a **commit log**
- ❑ Recently committed updates are stored in memory –**MEMtable**
- ❑ Older updates are stored in a sequence of **SSTables**.



Tablet Serving

❑ Write operation:

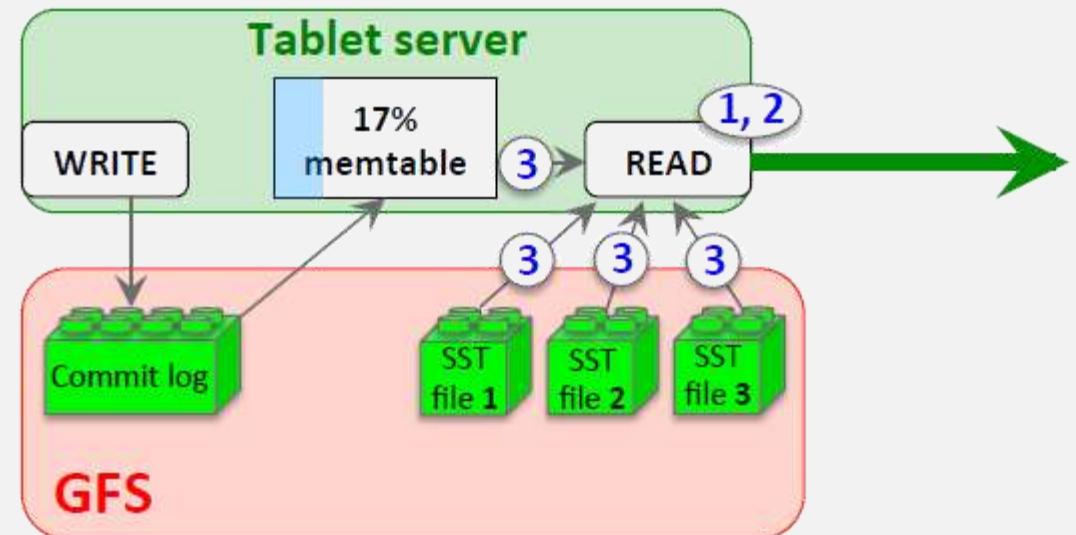
1. Server checks that the request is **well-formed**
2. Server checks that the sender is **authorized** to write (list of permitted writers in a Chubby file)
3. A **valid mutation** is written to the commit log that stores redo records (group commit to improve throughput)
4. After the mutation has been committed, its **contents** are inserted into the **MEMtable** (= in memory sorted buffer)



Tablet Serving

□ Read operation:

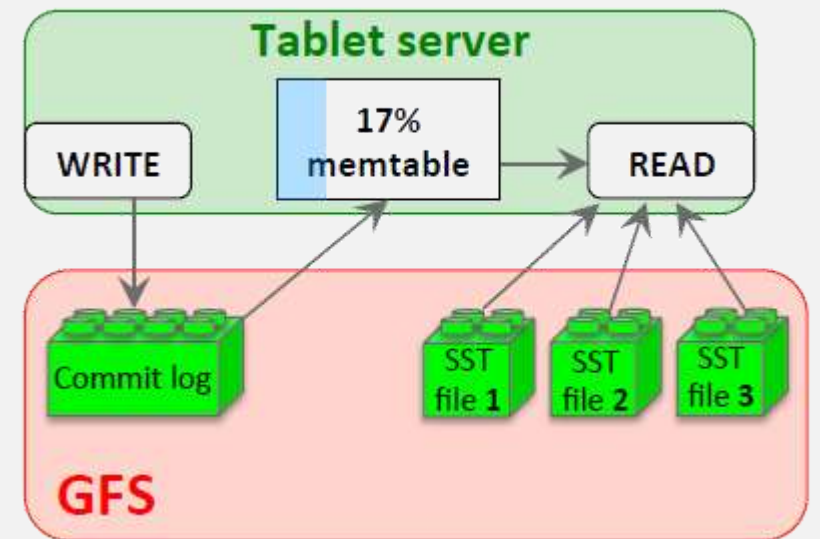
1. Server checks that the request is **well-formed**
2. Server checks that the sender is **authorized** to read (list of permitted writers in a Chubby file)
3. **Valid read operation** is executed on a merged view of the sequence of SSTables and the **MEMtable**



Tablet Serving

❑ Tablet Recovery

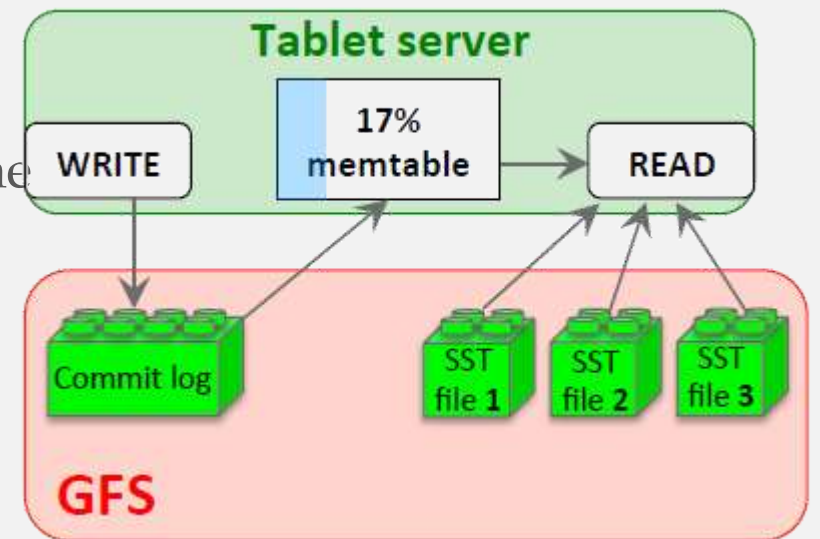
1. Tablet server reads its metadata from the METADATA table (lists of SSTables that comprise a tablet and a set of redo points, which are pointers into any commit logs that may contain data for the tablet)
2. The tablet server reads the indices of the SSTables into memory and reconstructs the MEMtable by applying all of the updates that have a committed since the redo points



Compaction

- In order to control size of **MEMtable**, **tablet log**, and **SSTable** files, “compaction” is used.


1. **Minor Compaction** - Move data from **MEMtable** to **SSTable**.
2. **Merging Compaction** - Merge multiple **SSTables** and **MEMtable** to a single **SSTable**.
3. **Major Compaction** - that re-writes **all** **SSTables** into exactly one **SSTable**



Compaction

1. Minor Compaction

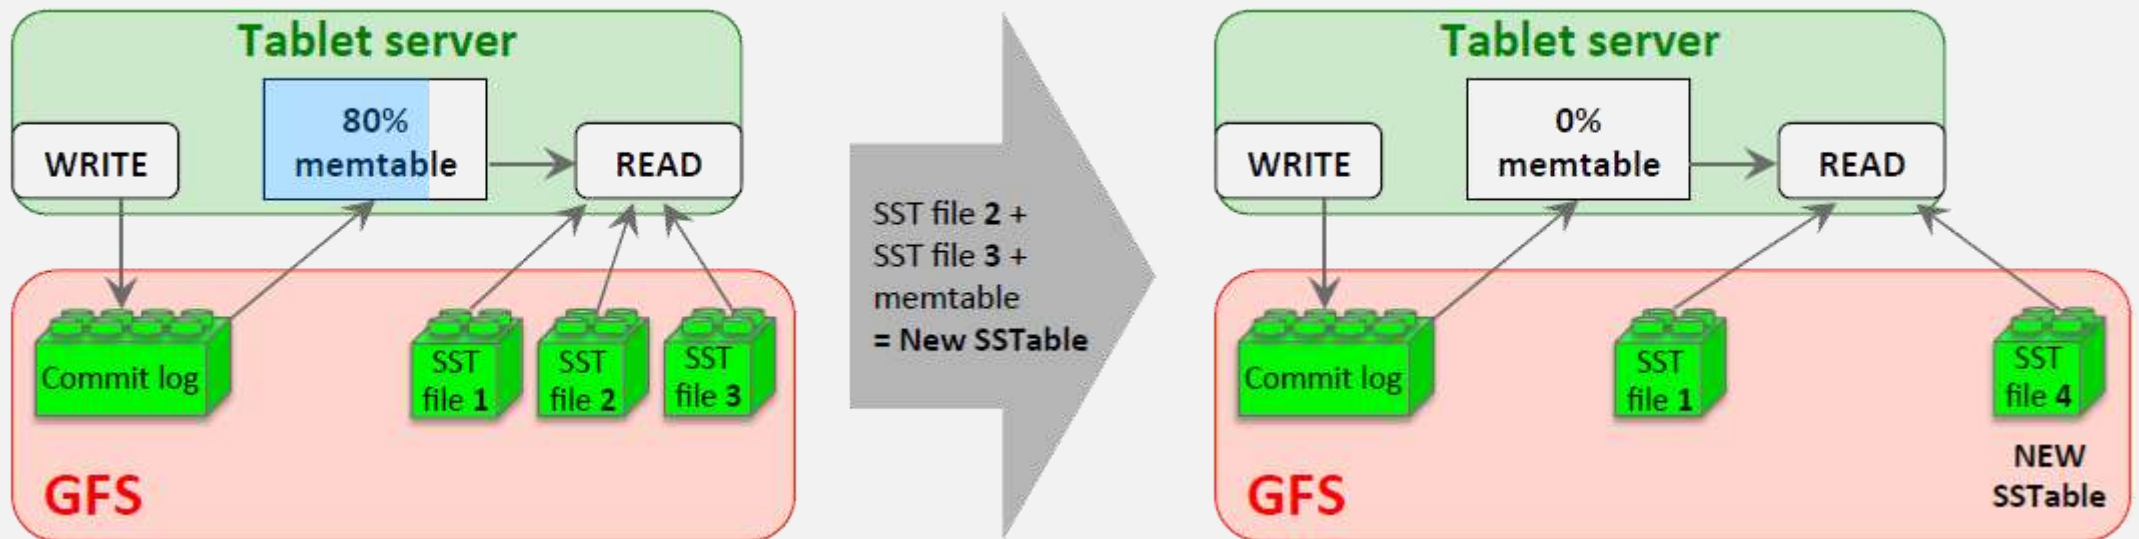
- When MEMtable size reaches a threshold, MEMtable is frozen, a new MEMtable is created, and the frozen MEMtable is converted to a new SSTable and written to GFS
- Two goals: shrinks the memory usage of the tablet server, reduces the amount of data that has to be read from the commit log during a recovery



Compaction

2. Merging Compaction

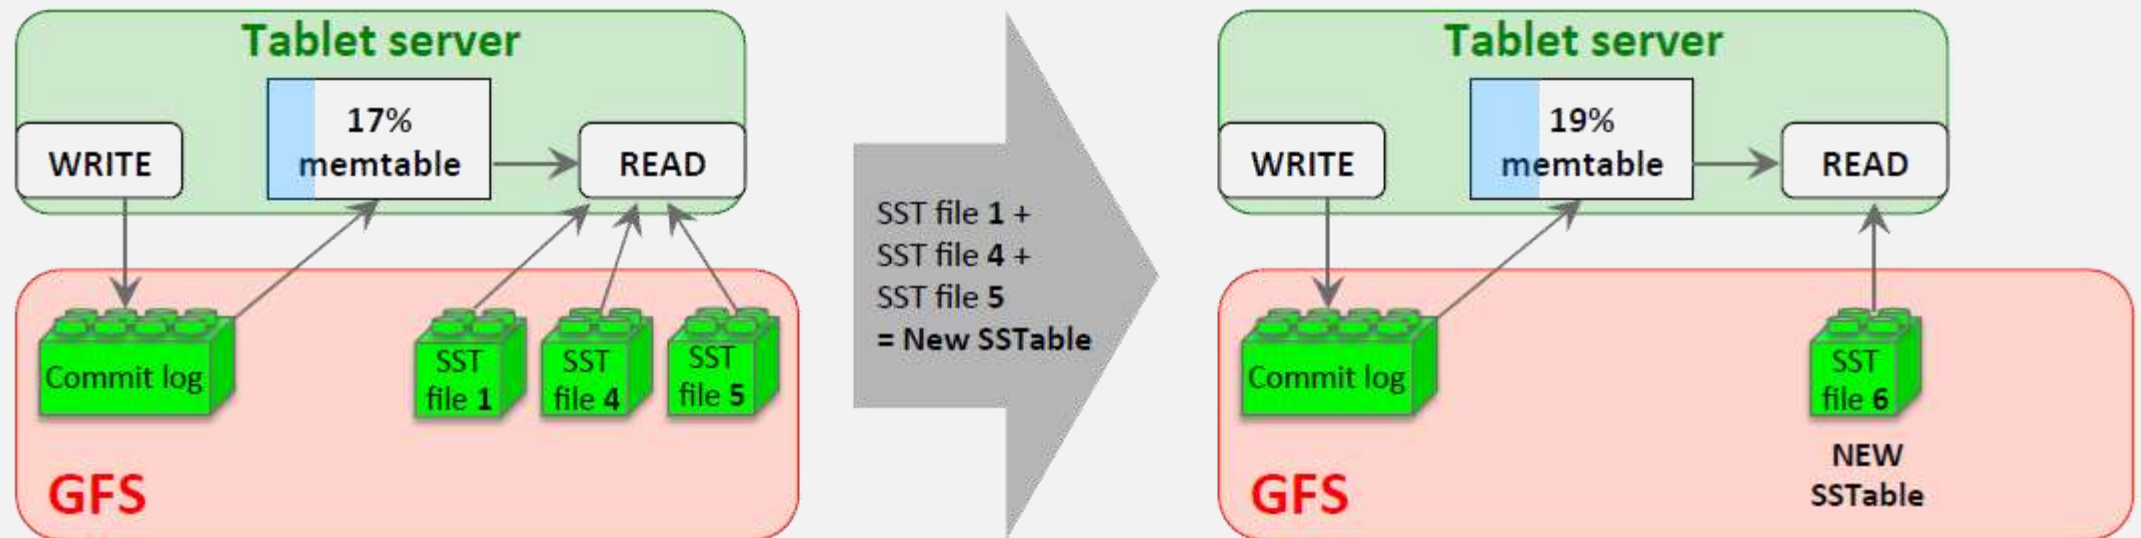
- **Problem:** every minor compaction creates a **new** SSTable (→ arbitrary number of SSTables !)
- **Solution:** periodic merging of a few SSTables and the MEMtable



Compaction

3. Major Compaction

- It is a merging compaction that rewrites all SSTables into exactly **one SSTable that contains no deletion information or deleted data**
- BigTable cycles through all of its tablets and regularly applies major compaction to them (=reclaim resources used by deleted data in a timely fashion)



Conclusion

- ❑ BigTable has achieved its goals of high performance, data availability and scalability.
- ❑ It has been successfully deployed in real apps (Personalized Search, Orkut, Google Maps, ...)
- ❑ Significant advantages of building own storage system like flexibility in designing data model, control over implementation and other infrastructure on which Bigtable relies on.



Thanks For Listening

