




# Python and Big data - An Introduction to Spark (PySpark)

Hitesh Dharmdasani

# About me

- Security Researcher, Malware Reversing Engineer, Developer
- GIT > GMU > Berkeley > FireEye > On Stage
- Bootstrapping a few ideas
- Hiring!



# What we will talk about?

- What is Spark?
- How does spark do things
- PySpark and data processing primitives
- Example Demo - Playing with Network Logs
- Streaming and Machine Learning in Spark
- When to use Spark

<http://bit.do/PyBelgaumSpark>

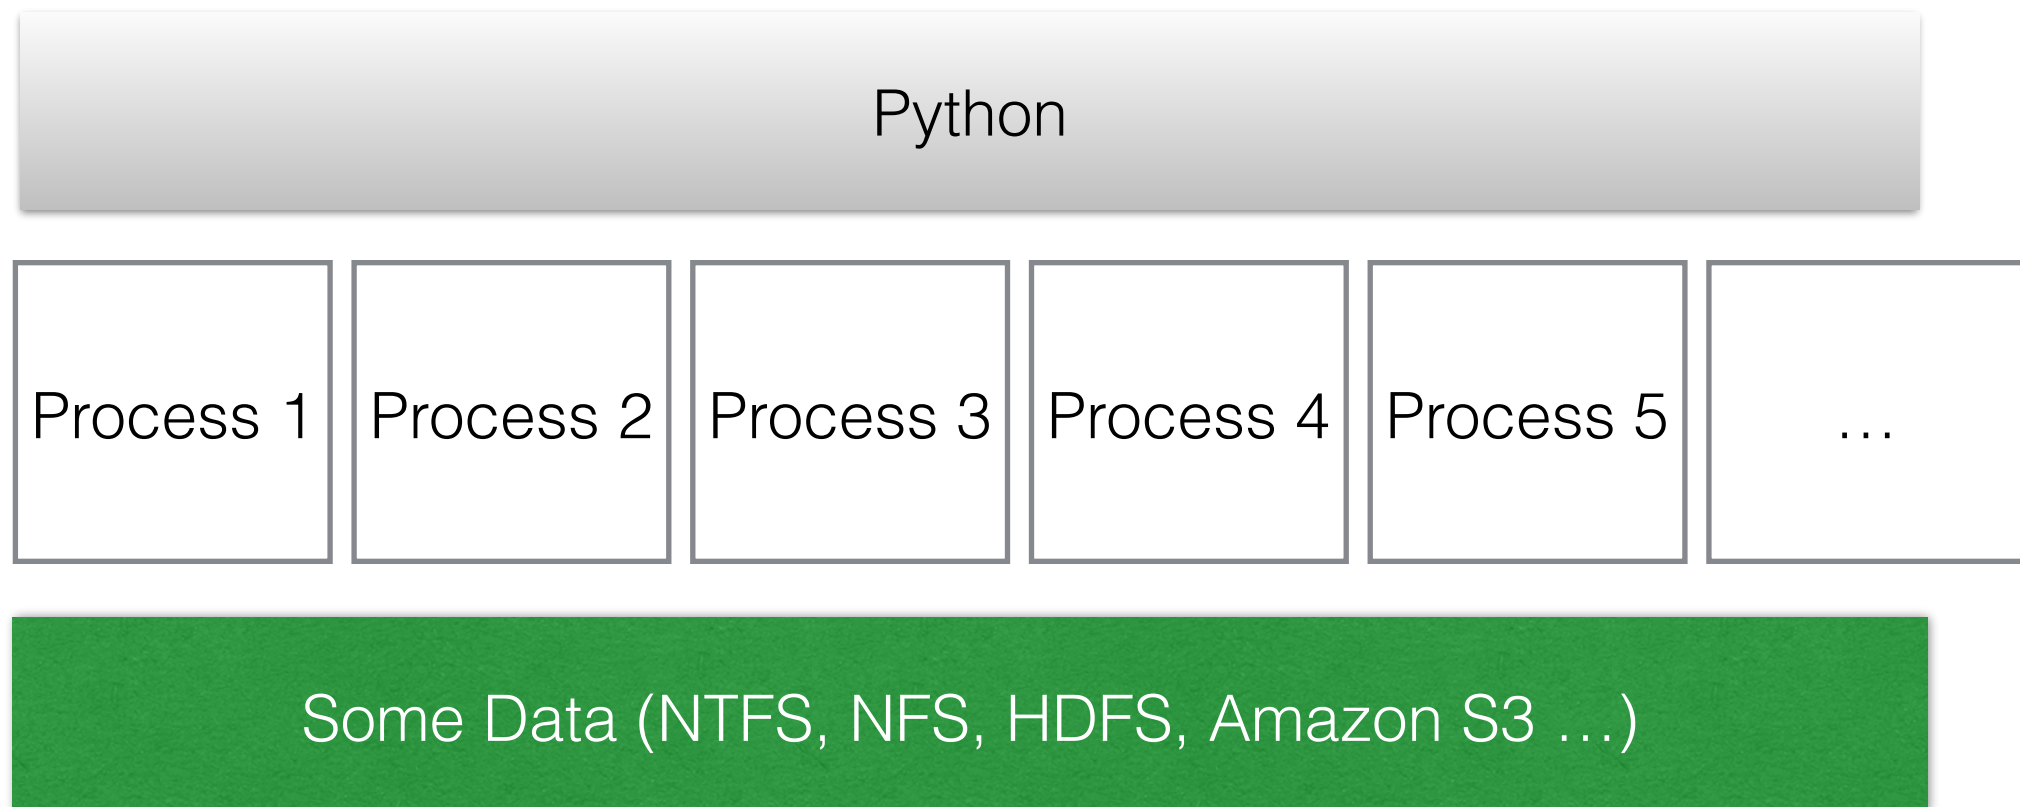
<http://tinyurl.com/PyBelgaumSpark>

# What will we NOT talk about

- Writing production level jobs
- Fine Tuning Spark
- Integrating Spark with Kafka and the like
- Nooks and Crooks of Spark
- But glad to talk about it offline

# The Common Scenario

You write 1 job. Then chunk, cut, slice and dice



# Compute where the data is

- Paradigm shift in computing
- Don't load all the data into one place and do operations
- State your operations and send code to the machine
- Sending code to machine >>> Getting data over network

# MapReduce

```
public static MyFirstMapper {  
    public void map { . . . }  
}
```

```
public static MyFirstReducer {  
    public void reduce { . . . }  
}
```

```
public static MySecondMapper {  
    public void map { . . . }  
}
```

```
public static MySecondReducer {  
    public void reduce { . . . }  
}
```

```
Job job = new Job(conf,  
    "First");  
job.setMapperClass(MyFirstMapper  
    .class);  
job.setReducerClass(MyFirstReduc  
    er.class);
```

```
/*Job 1 goes to Disk */
```

```
if(job.isSuccessful()) {
```

```
    Job job2 = new  
    Job(conf, "Second");  
    job2.setMapperClass(MySecondMap  
        per.class);  
    job2.setReducerClass(MySecondRe  
        ducer.class);
```

```
}
```

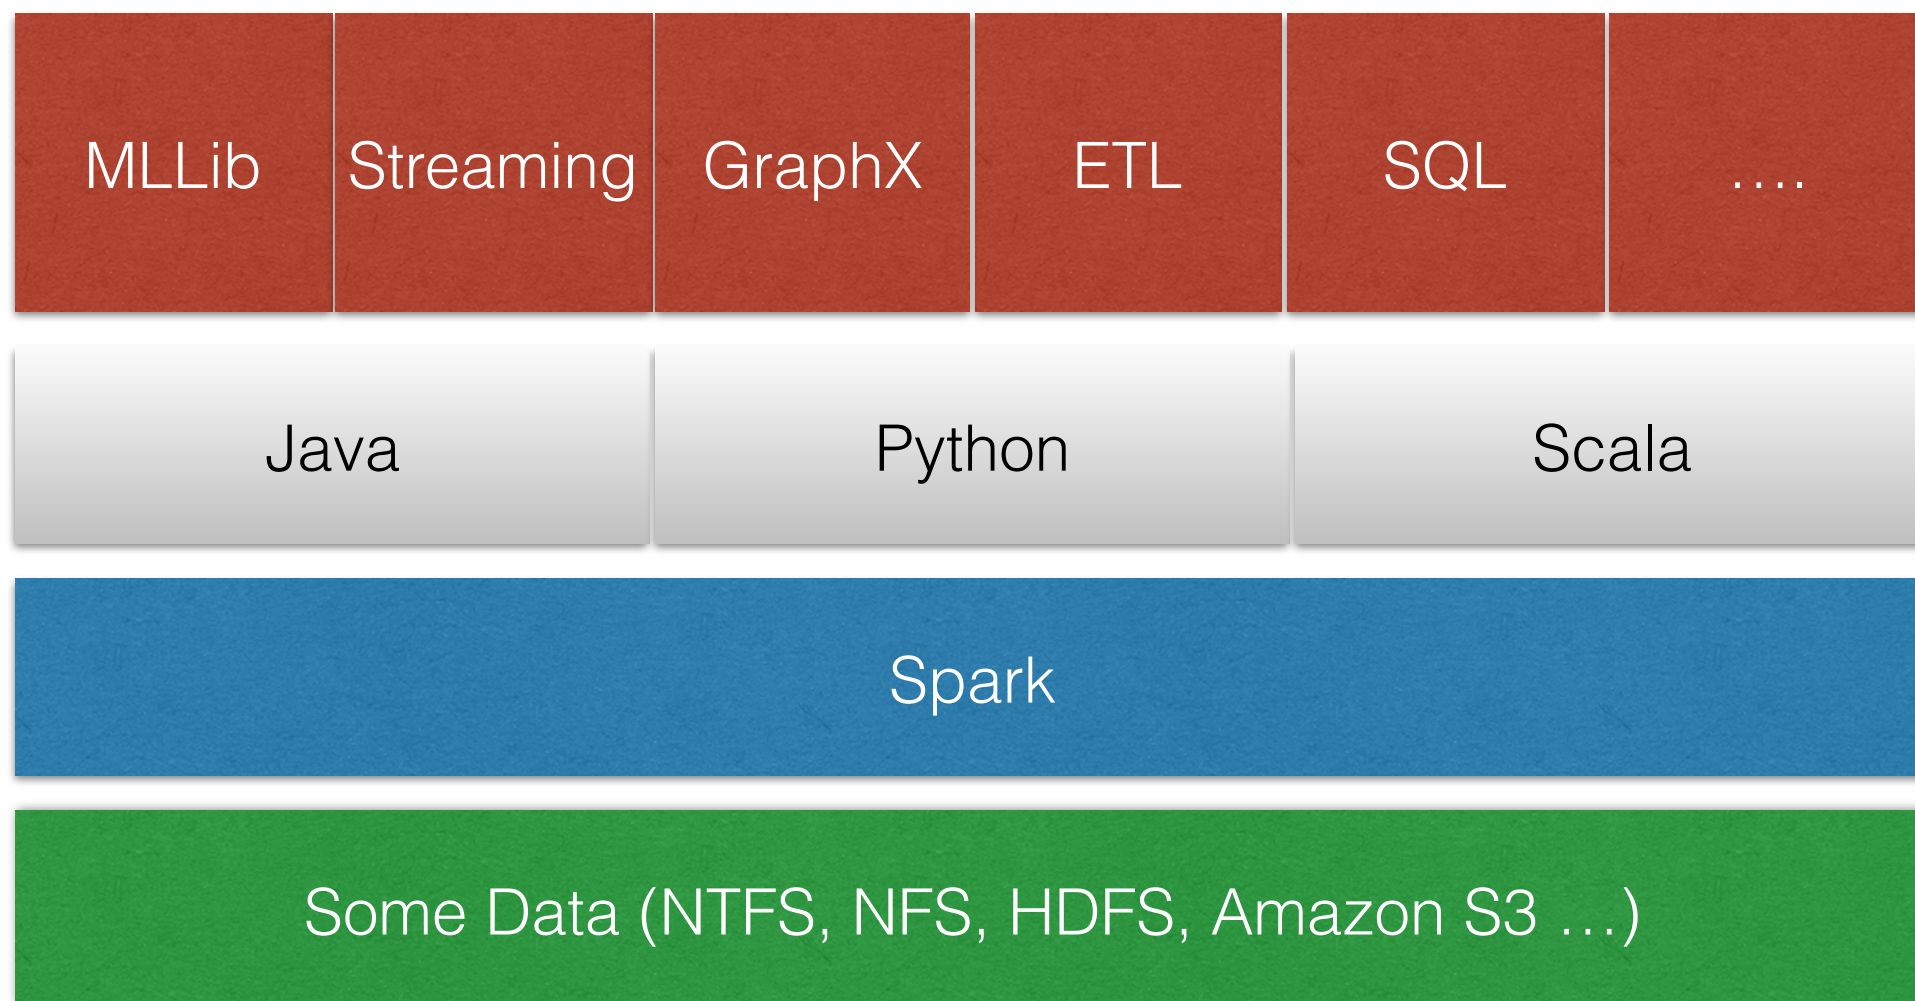
This also looks ugly if you ask me!

# What is Spark?

- Open Source Lightning Fast Cluster Computing
- Focus on Speed and Scale
- Developed at AMP Lab, UC Berkeley by Matei Zaharia
- Most active Apache Project in 2014 (Even more than Hadoop)
- **Recently beat MapReduce in sorting 100TB of data by being 3X faster and using 10X fewer machines**



# What is Spark?







# What is Spark?

- Inherently distributed
  - Computation happens where the data resides

Spark

Some Data (NTFS, NFS, HDFS, Amazon S3 ...)



# What is different from MapReduce

- Uses main memory for caching
- Dataset is partitioned and stored in RAM/Disk for iterative queries
- Large speedups for iterative operations when in-memory caching is used

# Spark Internals

## The Init

- Creating a SparkContext
- It is Sparks' gateway to access the cluster
- In interactive mode. SparkContext is created as 'sc'

```
$ pyspark
```

```
...
```

```
...
```

```
SparkContext available as sc.
```

```
>>> sc
```

```
<pyspark.context.SparkContext at 0xdeadbeef>
```

# Spark Internals

## The Key Idea

### Resilient Distributed Datasets

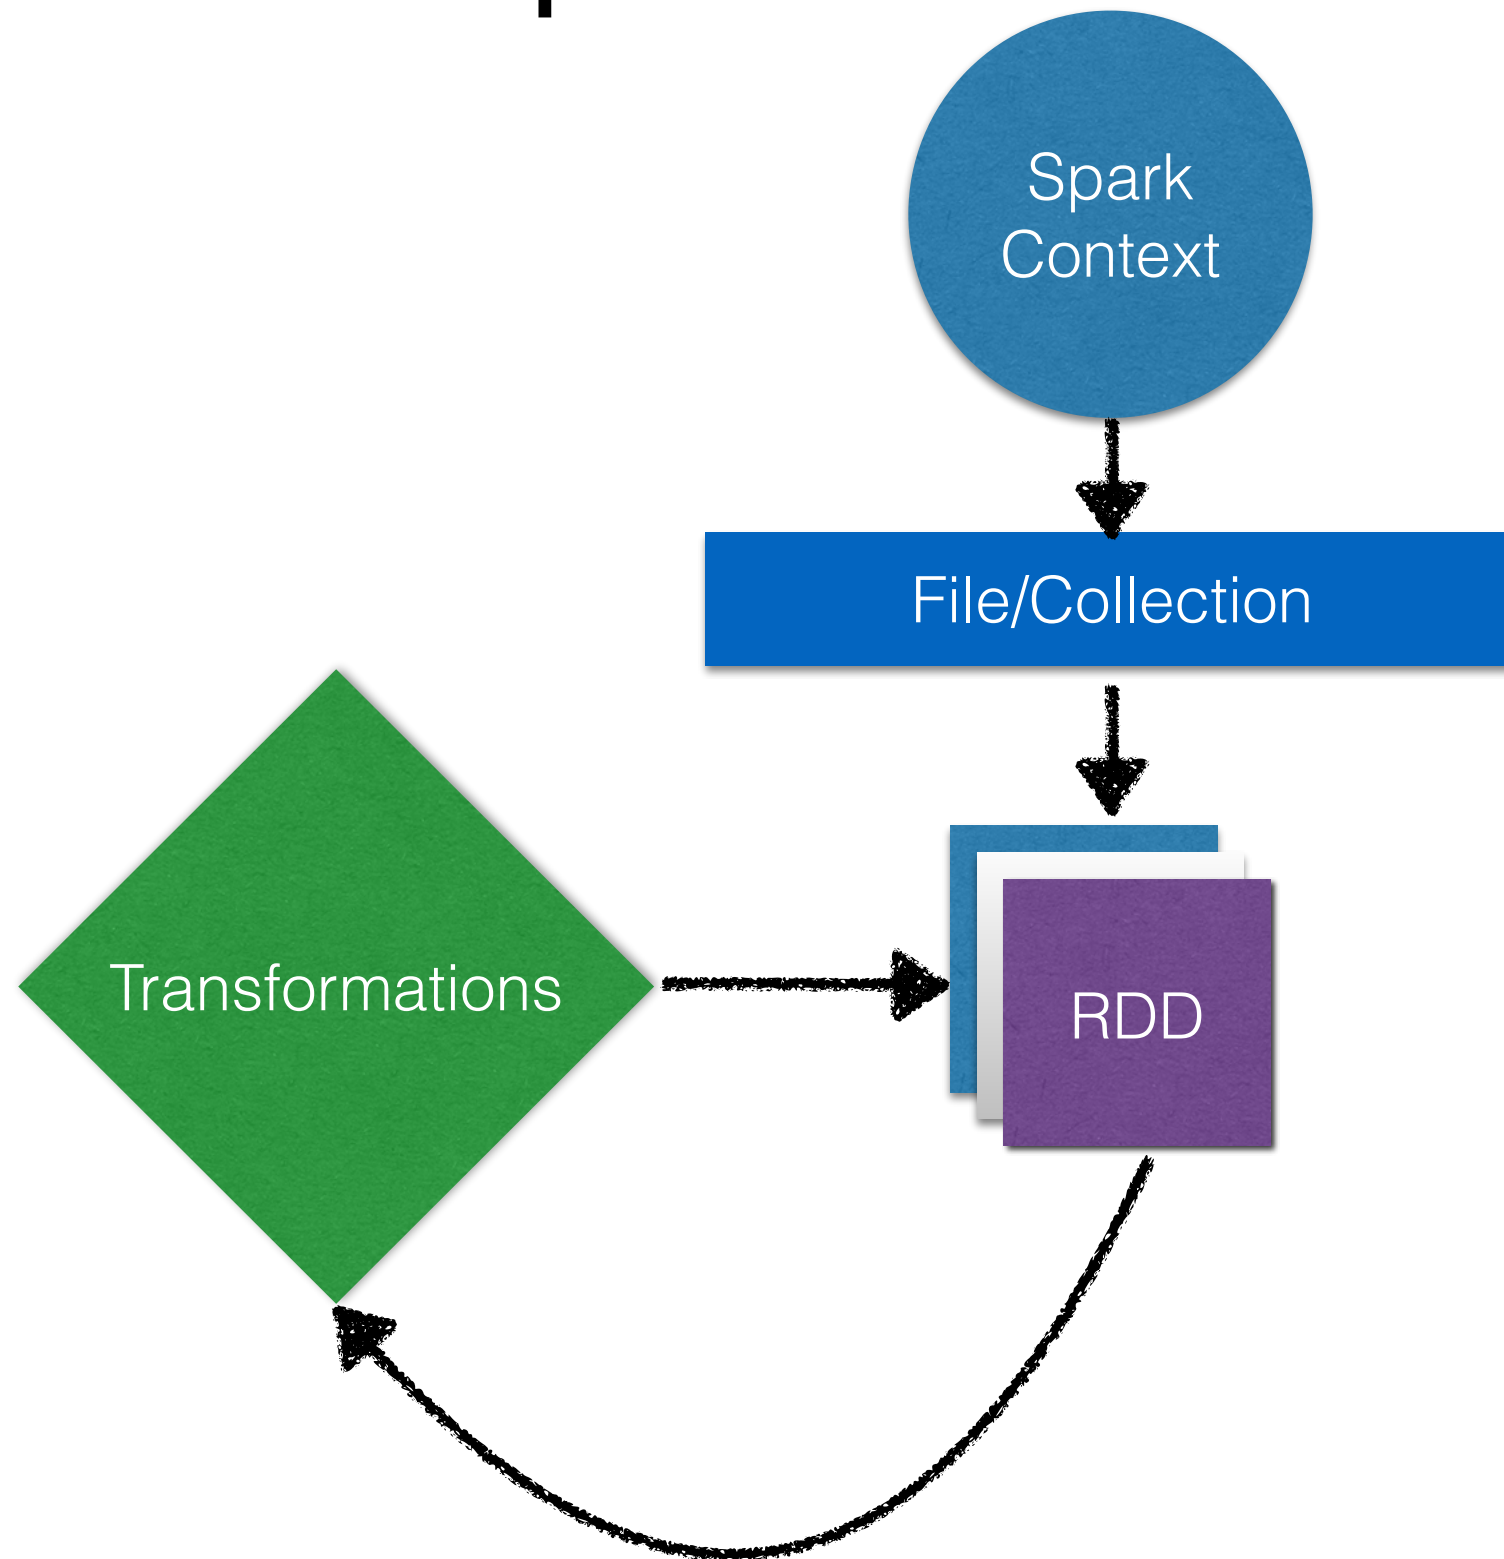
- Basic unit of abstraction of data
- Immutable
- Persistence

```
>>> data = [90, 14, 20, 86, 43, 55, 30, 94 ]  
>>> distData = sc.parallelize(data)  
ParallelCollectionRDD[13] at parallelize at  
PythonRDD.scala:364
```

# Spark Internals

Operations on RDDs - Transformations & Actions

# Spark Internals



# Spark Internals

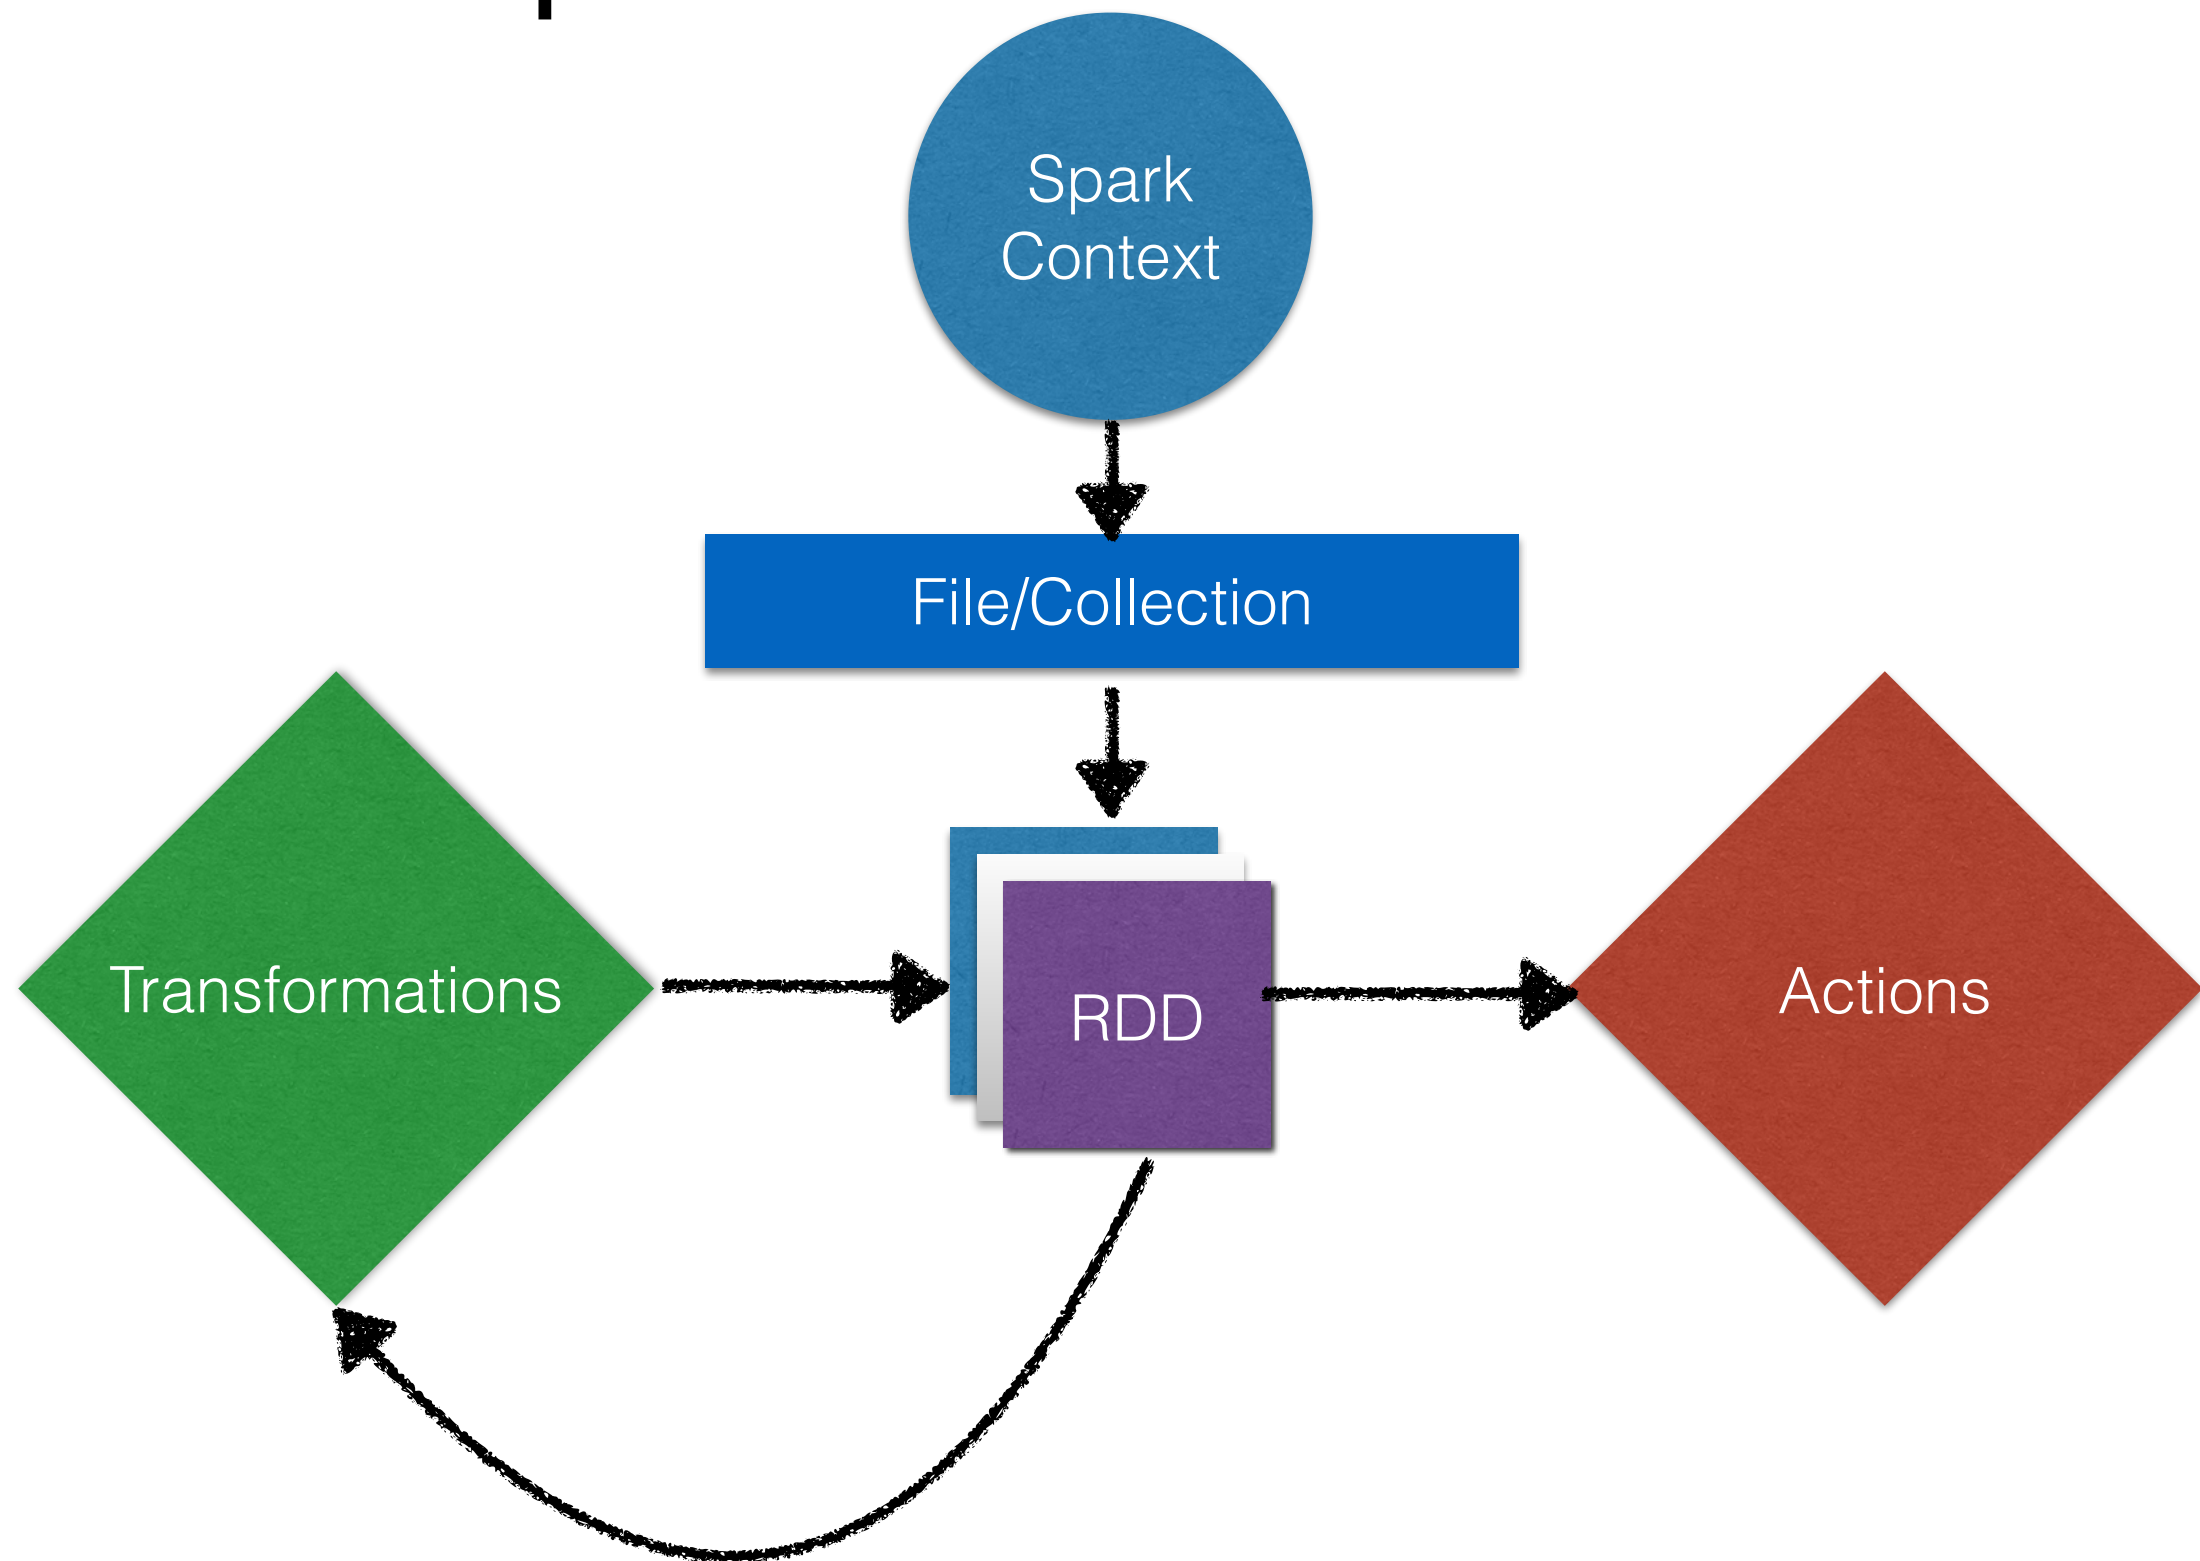
Now what?

Lazy Evaluation





# Spark Internals



# Spark Internals

## Transformation Operations on RDDs

### Map

```
def mapFunc(x):  
    return x+1
```

```
rdd_2 = rdd_1.map(mapFunc)
```

---

### Filter

```
def filterFunc(x):  
    if x % 2 == 0:  
        return True  
    else:  
        return False
```

```
rdd_2 = rdd_1.filter(filterFunc)
```

# Spark Internals

## Transformation Operations on RDDs

- `map`
- `filter`
- `flatMap`
- `mapPartitions`
- `mapPartitionsWithIndex`
- `sample`
- `union`
- `intersection`
- `distinct`
- `groupByKey`

# Spark Internals

```
>>> increment_rdd = distData.map(mapFunc)
>>> increment_rdd.collect()
[91, 15, 21, 87, 44, 56, 31, 95]
>>>
```

```
>>> increment_rdd.filter(filterFunc).collect()
[44, 56]
```

OR

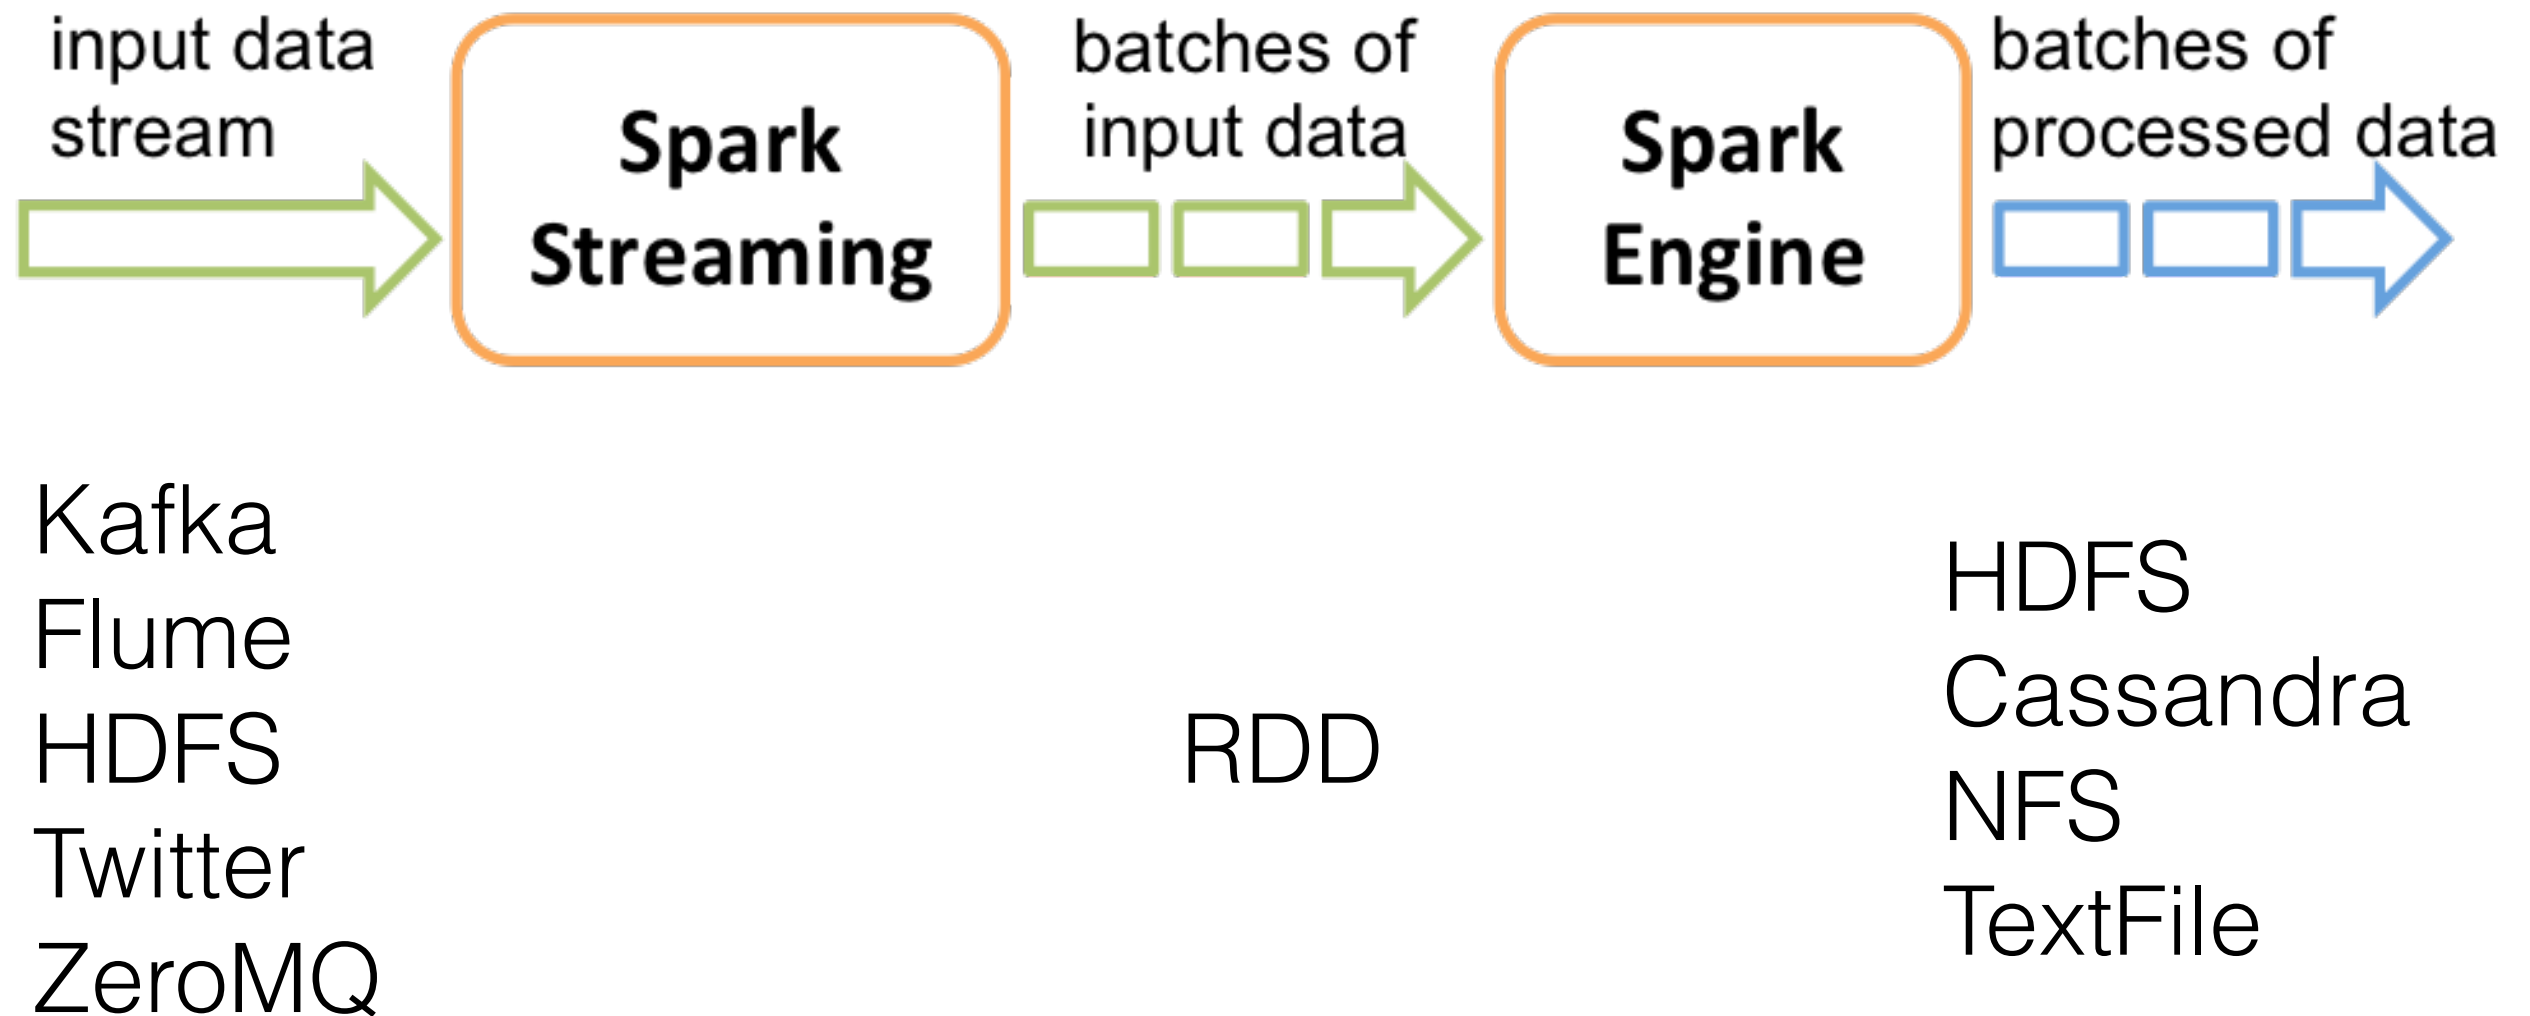
```
>>> distData.map(mapFunc).filter(filterFunc).collect()
[44, 56]
```

# Spark Internals

Fault Tolerance and Lineage

# Moving to the Terminal

# Spark Streaming

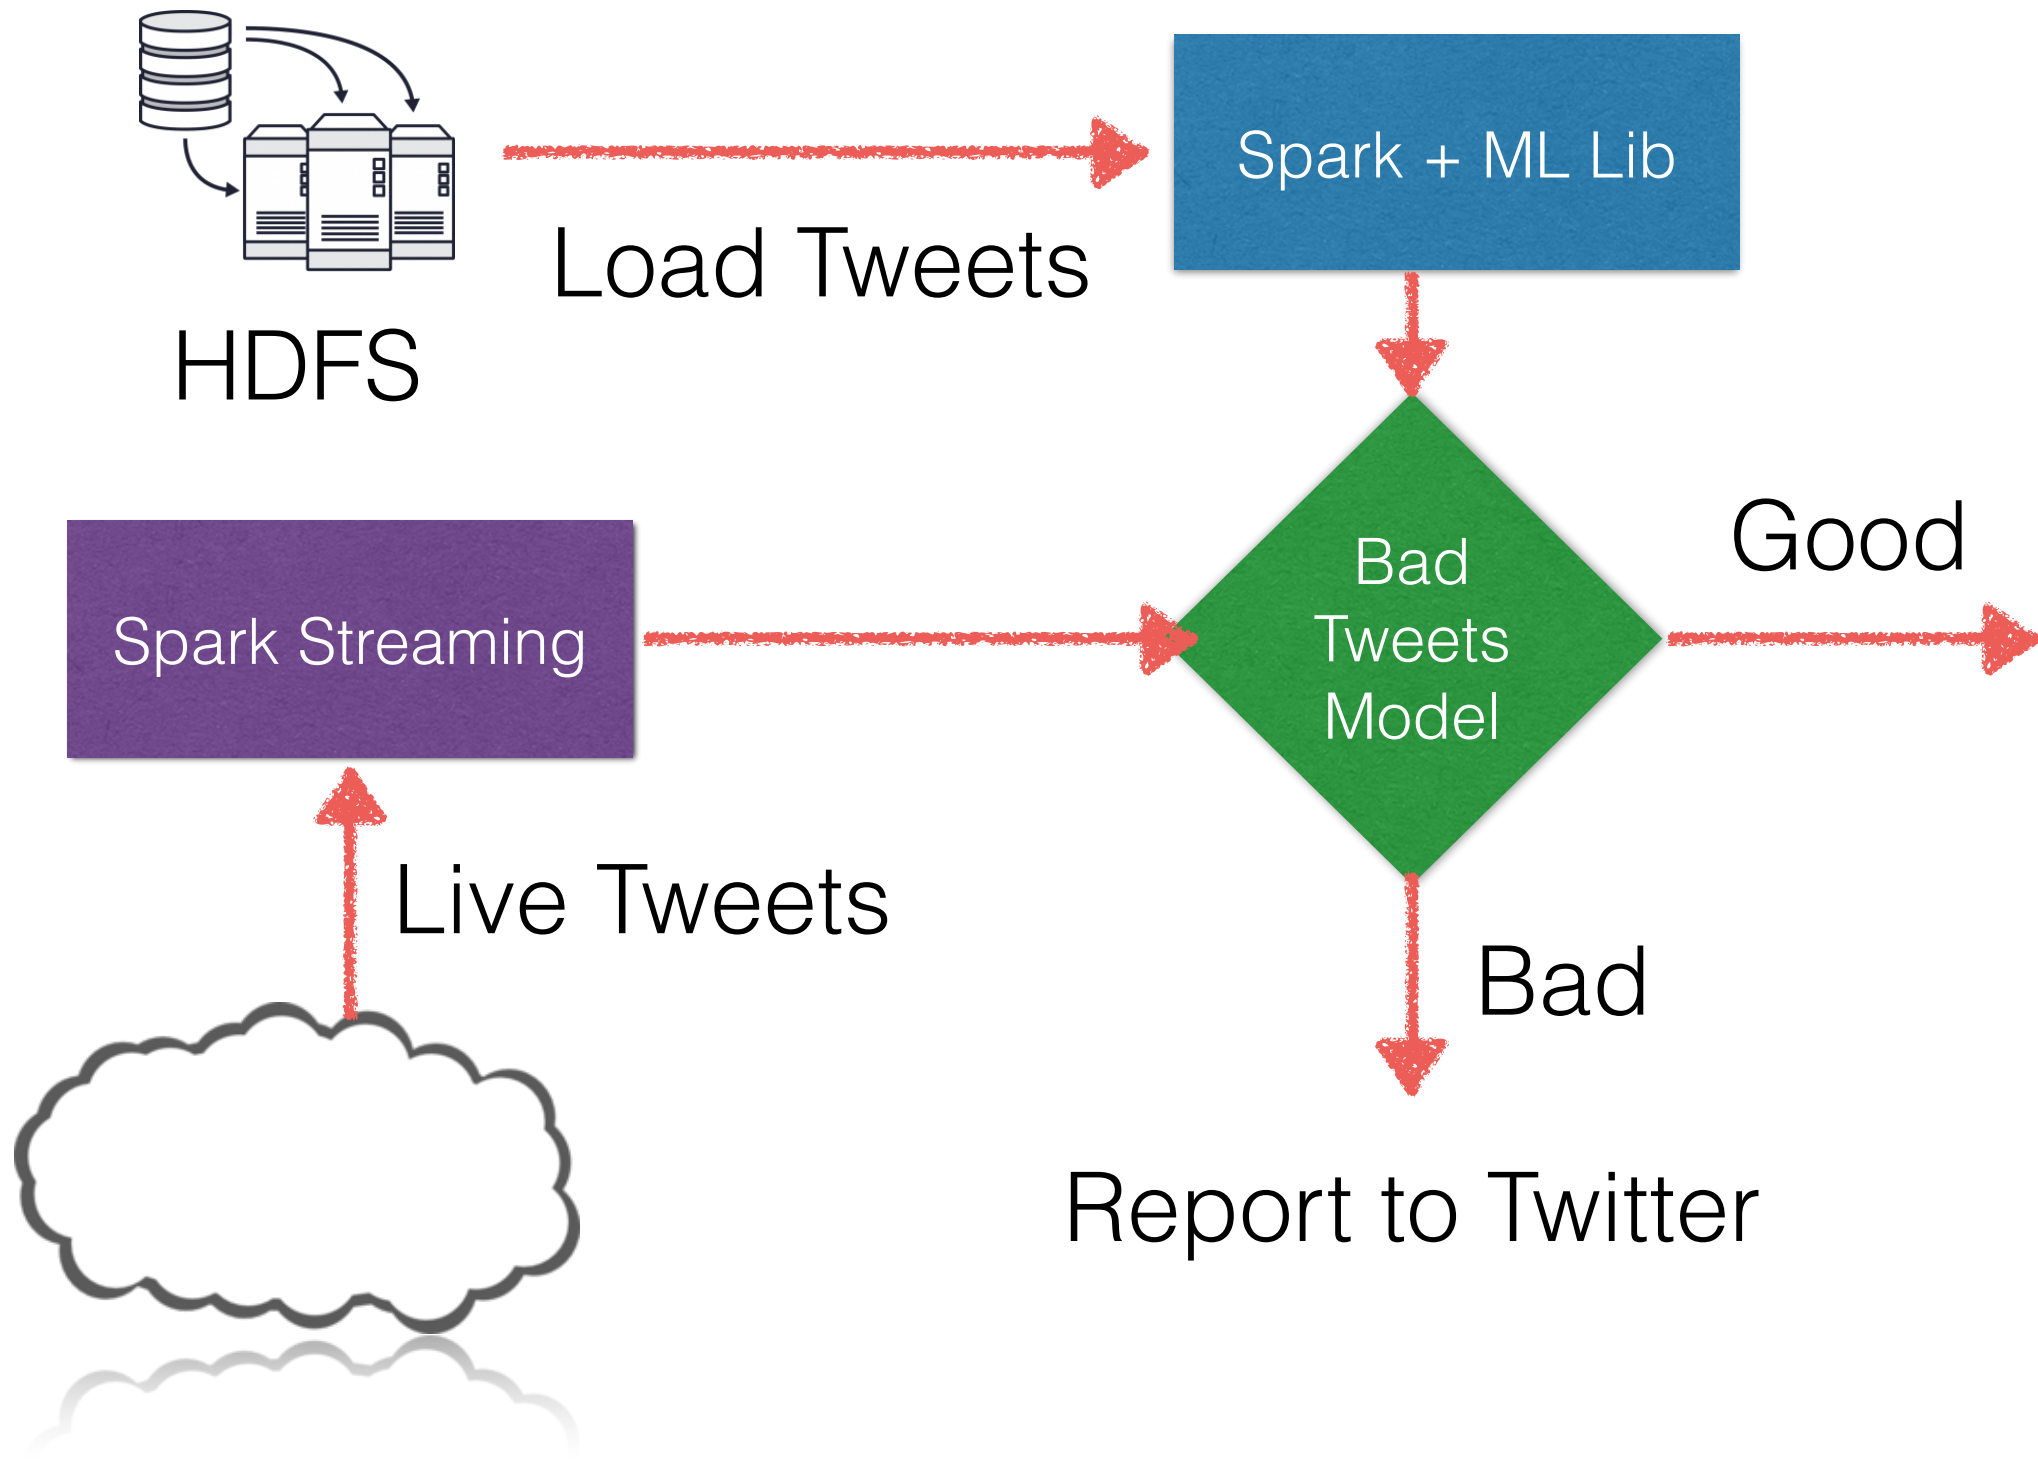


# ML Lib

- Machine Learning Primitives in Spark
- Provides training and classification at scale
- Exploits Sparks' ability for iterative computation (Linear Regression, Random Forest)
- Currently the most active area of work within Spark



# How can I use all this?



# To Spark or not to Spark

- Iterative computations
- “Don't fix something that is not broken”
- Lesser learning barrier
- Large one-time compute
- Single Map Reduce Operation